

# predicting-diabetes07-06

June 7, 2024

## 1 PREDICTING DIABETES IN PIMA INDIAN WOMEN USING MACHINE LEARNING MODELS

Capstone Project-I submitted to Imarticus Learning

Submitted By Dr. DILEEP KUMAR SHETTY

A population of women who were at least 21 years old, of Pima Indian heritage and living near Phoenix, Arizona, was tested for diabetes according to World Health Organization criteria. The data were collected by the US National Institute of Diabetes and Digestive and Kidney Diseases. Some values are not within the expected range and should be treated as missing values. What method would be more effective for filling this type of missing value? Additionally, how should further classification proceed? npreg: Number of times pregnant

glu: Plasma glucose concentration a 2 hours in an oral glucose tolerance test

bp: Diastolic blood pressure (mm Hg)

skin:Skin Thickness(Triceps skin fold thickness (mm))

insulin: 2-Hour serum insulin ( $\mu$ U/ml)

bmi: Body mass index ( $\text{weight in kg}/(\text{height in m})^2$ )

ped: Diabetes pedigree function

age: Age (years)

type –yes or no for diabetic according to WHO criteria

## 2 OBJECTIVES

This project aims to examine the factors influencing diabetes among Pima Indian women. The main objectives are as follows:

- i. To analyze the data on Pima Indian women using various ML models.
- ii. To explore whether specific characteristics or factors within the studied population are linked to a higher likelihood of having diabetes. Identifying such sub-groups is valuable for understanding risk factors and tailoring interventions or treatments more effectively.

### 3 1.1 Libraries and modules commonly used in data analysis and machine learning in Python

```
[1]: #Pandas is a powerful data manipulation library for Python.
import pandas as pd

#NumPy is a numerical computing library for Python.
import numpy as np

#Matplotlib is a plotting library for creating static, interactive, and
↳ animated visualizations in Python.
import matplotlib.pyplot as plt

#ListedColormap is a class in Matplotlib used to create a colormap from a list
↳ of colors.
from matplotlib.colors import ListedColormap

#Seaborn is a statistical data visualization library based on Matplotlib.
import seaborn as sns

#is_string_dtype is a function from Pandas used to check if a dtype is of
↳ object type.
from pandas.api.types import is_string_dtype

#StandardScaler is a preprocessing technique used to standardize features by
↳ removing the mean and scaling to unit variance.
from sklearn.preprocessing import StandardScaler

#train_test_split is a function in scikit-learn used for splitting a dataset
↳ into training and testing sets.
from sklearn.model_selection import train_test_split

[2]: #The metrics module in scikit-learn provides various metrics for evaluating
↳ model performance.
from sklearn import metrics

#LogisticRegression is a class in scikit-learn used for logistic regression
↳ modeling.
from sklearn.linear_model import LogisticRegression

#classification_report is a function in scikit-learn that generates a text
↳ report showing the main classification metrics.
from sklearn.metrics import classification_report

#cohen_kappa_score is a function in scikit-learn used for calculating the
↳ Cohen's kappa statistic.
```

```

from sklearn.metrics import cohen_kappa_score

#confusion_matrix is a function in scikit-learn that computes the confusion
↳matrix to evaluate the accuracy of a classification.
from sklearn.metrics import confusion_matrix

#roc_auc_score is a function in scikit-learn used for computing the area under
↳the ROC AUC.
from sklearn.metrics import roc_auc_score

#roc_curve is a function in scikit-learn used for generating receiver operating
↳characteristic (ROC) curves.
from sklearn.metrics import roc_curve

#SGDClassifier is a class in scikit-learn implementing linear classifiers with
↳Stochastic Gradient Descent training.
from sklearn.linear_model import SGDClassifier

#DecisionTreeClassifier is a class in scikit-learn for building decision tree
↳models.
from sklearn.tree import DecisionTreeClassifier

#GridSearchCV is a class in scikit-learn for hyperparameter tuning using grid
↳search.
from sklearn.model_selection import GridSearchCV

#The tree module in scikit-learn provides tools for working with decision trees.
from sklearn import tree

#export_graphviz is a function in scikit-learn for exporting decision tree
↳models to Graphviz format.
from sklearn.tree import export_graphviz

```

```

[3]: #Statsmodels is a library for estimating and testing statistical models.
import statsmodels
import statsmodels.api as sm

#SVC is a class in scikit-learn implementing Support Vector Classification.
from sklearn.svm import SVC

#GaussianNB is a class in scikit-learn implementing Gaussian Naive Bayes
↳classification.
from sklearn.naive_bayes import GaussianNB

#KNeighborsClassifier is a class in scikit-learn for k-nearest neighbors
↳classification.

```

```
[4]: #Ignore Warnings:
import warnings
from warnings import filterwarnings
filterwarnings('ignore')

#Adjust Figure Size for Matplotlib:
plt.rcParams['figure.figsize'] = [10,4]

[5]: #Adjusting some display and print options for Pandas and NumPy
#max_columns to None, Pandas not to truncate the display of columns.
pd.options.display.max_columns = None

##max_rows to None, Pandas not to truncate the display of rows.
pd.options.display.max_rows = None

# To see the full numeric values without exponential notation.
np.set_printoptions(suppress=True)

[6]: #The os.chdir function is used to change the current working directory to the
      ↪specified path.
import os
os.chdir("C:\DKS\Machine_Learning\Extra_Projects")

##Load the Dataset
data = pd.read_csv('diabetes.csv')

#The sample(15) method is used to display a random sample of 15 rows from the
      ↪loaded DataFrame
data.sample(15)
```

```
[6]:      npreg  glu  bp  skin  insulin   bmi    ped  age  type
25         10  125  70   26      115  31.1  0.205  41    1
152         9  156  86   28      155  34.3  1.189  42    1
466         0   74  52   10       36  27.8  0.269  22    0
323        13  152  90   33       29  26.8  0.731  43    1
465         0  124  56   13      105  21.8  0.452  21    0
603         7  150  78   29      126  35.2  0.692  54    1
172         2   87   0   23        0  28.9  0.773  25    0
372         0   84  64   22       66  35.8  0.545  21    0
320         4  129  60   12      231  27.5  0.527  31    0
233         4  122  68    0        0  35.0  0.394  29    0
701         6  125  78   31        0  27.6  0.565  49    1
585         1   93  56   11        0  22.5  0.417  22    0
750         4  136  70    0        0  31.2  1.182  22    1
397         0  131  66   40        0  34.3  0.196  22    1
755         1  128  88   39      110  36.5  1.057  37    1
```

```
[7]: # Display summary statistics
summary_stats = data.describe()
summary_stats
```

```
[7]:
```

	npreg	glu	bp	skin	insulin	bmi \
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000

	ped	age	type
count	768.000000	768.000000	768.000000
mean	0.471876	33.240885	0.348958
std	0.331329	11.760232	0.476951
min	0.078000	21.000000	0.000000
25%	0.243750	24.000000	0.000000
50%	0.372500	29.000000	0.000000
75%	0.626250	41.000000	1.000000
max	2.420000	81.000000	1.000000

In the variables Glucose, BP, Skin, Insulin, and BMI, the minimum values are zero. In reality, these factors are unlikely to have a minimum value of zero; instead, these zero values are indicative of missing data. Therefore, we need to replace these zeros with NA to accurately represent the absence of valid information.

```
[8]: #This lines selects specific columns from the original DataFrame (data)
data_Correct=data.iloc[:,[0,6,7,8]]
data_missing=data.iloc[:,1:6]

#This line replaces all occurrences of 0 in the data_missing DataFrame with NaN.
↪
#The inplace=True argument modifies the DataFrame in place.
data_missing.replace(0, np.nan, inplace=True)

#Concatenates (pd.concat) the two DataFrames (data_Correct and data_missing)
↪along the columns (axis=1).
#The result is stored in the variable data.
data= pd.concat([data_Correct,data_missing], axis=1)
data.head()
```

```
[8]:
```

	npreg	ped	age	type	glu	bp	skin	insulin	bmi
0	6	0.627	50	1	148.0	72.0	35.0	NaN	33.6
1	1	0.351	31	0	85.0	66.0	29.0	NaN	26.6

2	8	0.672	32	1	183.0	64.0	NaN	NaN	23.3
3	1	0.167	21	0	89.0	66.0	23.0	94.0	28.1
4	0	2.288	33	1	137.0	40.0	35.0	168.0	43.1

```
[9]: #The dtypes attribute in Pandas is used to display the data types of each
      ↪column in a DataFrame.
data.dtypes
```

```
[9]: npreg      int64
ped          float64
age          int64
type         int64
glu          float64
bp           float64
skin         float64
insulin      float64
bmi          float64
dtype: object
```

```
[10]: # Check the info
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   npreg       768 non-null   int64
1   ped         768 non-null   float64
2   age         768 non-null   int64
3   type        768 non-null   int64
4   glu         763 non-null   float64
5   bp          733 non-null   float64
6   skin        541 non-null   float64
7   insulin     394 non-null   float64
8   bmi         757 non-null   float64
dtypes: float64(6), int64(3)
memory usage: 54.1 KB
```

```
[11]: #Splitting the DataFrame into feature variables (data_x) and the target
      ↪variable (data_y).
data_x = data.iloc[:, data.columns != 'type']
data_y = data.iloc[:,data.columns == 'type']
print(data_y.head(2))
print(data_x.head(2))
```

```
type
```

```

0      1
1      0
      npreg    ped    age    glu    bp    skin    insulin    bmi
0         6  0.627    50  148.0  72.0  35.0      NaN    33.6
1         1  0.351    31   85.0  66.0  29.0      NaN    26.6

```

```

[12]: #Summary statistics table where the features are listed as rows, and the
      ↳ summary statistics are columns.
      data_x.describe().T

```

```

[12]:      count      mean      std      min      25%      50%      75%  \
npreg    768.0    3.845052    3.369578    0.000    1.00000    3.0000    6.00000
ped       768.0    0.471876    0.331329    0.078    0.24375    0.3725    0.62625
age       768.0   33.240885   11.760232   21.000   24.00000   29.0000   41.00000
glu       763.0  121.686763   30.535641   44.000   99.00000  117.0000  141.00000
bp        733.0   72.405184   12.382158   24.000   64.00000   72.0000   80.00000
skin      541.0   29.153420   10.476982    7.000   22.00000   29.0000   36.00000
insulin   394.0  155.548223  118.775855   14.000   76.25000  125.0000  190.00000
bmi       757.0   32.457464    6.924988   18.200   27.50000   32.3000   36.60000

      max
npreg    17.00
ped        2.42
age       81.00
glu      199.00
bp       122.00
skin      99.00
insulin  846.00
bmi       67.10

```

## 4 Interpretation:

**Glucose Level:** The mean glucose level in the blood is approximately 119 mg/dL. It's worth noting that when fasting blood glucose falls between 100 to 125 mg/dL (5.6 to 6.9 mmol/L), it is typically recommended to consider lifestyle changes and monitor glycemia closely.

**Blood Pressure (BP):** The average blood pressure (BP) in our dataset is around 71 mmHg. It's important to mention that a normal blood pressure level is typically less than 120/80 mmHg.

**BMI (Body Mass Index):** The mean BMI in our dataset is 33.24. According to BMI categories, a BMI of 30.0 or higher falls within the obese range. Therefore, in our study, the average BMI suggests that participants are in the overweight range.

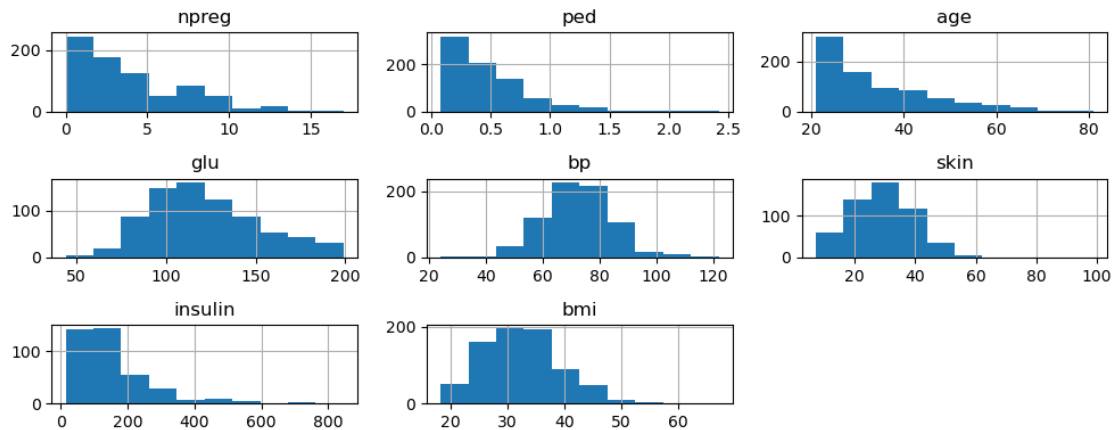
**Age:** The mean age in our study is 31 years, with the maximum age recorded being 67 years.

```

[13]: # plot the histogram of numeric features variables
      data_x.hist()
      plt.tight_layout()

```

```
plt.show()
```



Interpretation: In this dataset, the variables age, ped, and npreg follow right-skewed distributions. BP and BMI, on the other hand, exhibit approximately symmetric distributions, while the variable skin displays a bimodal distribution.

Conclusions: Age: The variable “age” in our dataset is characterized by a right-skewed distribution. This means that the majority of individuals in our sample tend to be younger, with a few outliers who are older than the average age.

Ped: Similarly, the variable “ped” also displays a right-skewed distribution. This implies that most observations have relatively low values, while a smaller proportion of the data have higher values for this variable.

Npreg: The variable “npreg” exhibits a right-skewed distribution as well. This suggests that the majority of individuals in our dataset have a low number of pregnancies, while a smaller subset of the data includes individuals with a higher number of pregnancies.

BP: In contrast to the aforementioned variables, “BP” follows an approximately symmetric distribution. This implies that the values of blood pressure (BP) are evenly distributed around the mean, with a similar number of individuals having both high and low blood pressure readings.

BMI: Like BP, the variable “BMI” also displays an approximately symmetric distribution. This suggests that the body mass index (BMI) values in our dataset are evenly spread around the mean, with a similar number of individuals having both low and high BMI values.


Skin: The variable “skin” is characterized by a bimodal distribution. This means that there are two distinct peaks or modes in the data, indicating that there are two groups of individuals in our sample with different values for this variable. These groups may have different characteristics related to skin measurements.

```
[14]: class_frequency = data_y.value_counts()  
      class_frequency
```



```
[14]: type
0      500
1      268
dtype: int64
```

```
[15]: class_frequency = data_y.value_counts()
class_frequency
sns.countplot(data=data_y, x="type")
plt.text(x = -0.05, y = data_y.value_counts()[0]+1, s =
↳str(round((class_frequency[0])*100/len(data_y),2)) + '%')
plt.text(x = 0.95, y = data_y.value_counts()[1]+1, s =
↳str(round((class_frequency[1])*100/len(data_y),2)) + '%')
plt.title('Count Plot for Target Variable Types', fontsize = 15)
plt.xlabel('Target Variable (Type)', fontsize = 15)
plt.ylabel('Count', fontsize = 15)
plt.show()
```



Interpretation: In our study, the target variable is “type,” which signifies whether a person has diabetes or not. Here, the value 0 denotes non-diabetic individuals, indicating that they do not have diabetes, while the value 1 represents diabetics, indicating that they have been diagnosed with diabetes. Our study reveals that 32.83% of the participants are classified as diabetics, while 67.17% are classified as non-diabetics.

Conclusion: In our research study, we have a target variable called “type,” which serves as an indicator of an individual’s diabetes status. This variable takes on two distinct values:

When “type” is assigned the value 0, it signifies that the individual is categorized as non-diabetic. In other words, individuals with this value do not have diabetes.


When “type” is assigned the value 1, it indicates that the individual is categorized as diabetic. This means they have been diagnosed with diabetes.

Based on our analysis of the data, we have found that 34.90% of the study participants fall into the diabetic category (type = 1), indicating that they have diabetes. In contrast, 65.1% of the participants fall into the non-diabetic category (type = 0), signifying that they do not have diabetes.

These percentages are important findings as they provide insights into the prevalence of diabetes within our study population. It suggests that a significant portion of the individuals in our study have been diagnosed with diabetes (34.90%), while the majority are non-diabetic (65.10%). This information is crucial for understanding the distribution of diabetes within the studied group and may have implications for healthcare interventions or further research related to diabetes management and prevention.


[ ]:

```
[16]: data_x.boxplot()
plt.title('Distribution of all Numeric Variables', fontsize = 15)
plt.xticks(rotation = 'vertical', fontsize = 15)
plt.show()
```



Notice that the variables 'npreg' and 'ped' has a quite small range as compared to the other variables. Thus, it is difficult to see the outliers for these variables. So, we plot the boxplot only for the variables 'npreg', 'ped'.

```
[17]: variables = [ 'npreg', 'ped' ]
data_x[variables].boxplot()
plt.title('Distribution of Independent Variables npreg and ped', fontsize = 15)
plt.xticks(rotation = 'vertical', fontsize = 15)
plt.show()
```



In our study, we've presented two plots that illustrate the distribution of data for various variables. These plots have revealed some interesting insights:

**Outliers in all the Variables:** For most of the variables we've analyzed, we've noticed the presence of data points that exceed the upper extreme of the distribution. This indicates that there are some values in our dataset that are significantly higher than the majority of the data. These extreme values are commonly referred to as outliers.

**Outliers in Blood Pressure (BP):** In the case of blood pressure (BP), we've identified a different scenario. Here, we have not only identified data points above the upper extreme but also some below the lower extreme. This suggests that there are both high and low extreme values in the BP variable.

Now, the decision we've made regarding these outliers is crucial and context-dependent. Given that our study is related to medical research, these extreme values may carry important clinical significance. Outliers in medical data can sometimes represent critical health conditions or unusual responses to treatment. Therefore, we have chosen not to remove these outliers from our dataset.

By retaining these extreme values, our analysis may capture the full spectrum of possible medical scenarios, including rare or unusual cases. This decision aligns with the goal of thorough medical research, which seeks to understand and account for all possible variations and conditions within the studied population.

```
[18]: Total = data.isnull().sum().sort_values(ascending = False)
Percentage = (data.isnull().sum()*100/data.isnull().count()).
↳sort_values(ascending = False)
Missing_Values = pd.concat([Total, Percentage], axis = 1, keys = ['Total',
↳'Percentage of missing observations'])
Missing_Values
```

```
[18]:
```

	Total	Percentage of missing observations
insulin	374	48.697917
skin	227	29.557292
bp	35	4.557292
bmi	11	1.432292
glu	5	0.651042
npreg	0	0.000000
ped	0	0.000000
age	0	0.000000
type	0	0.000000


Interpretation: With insulin having nearly 50% missing values in the dataset, it is prudent to consider removing this variable instead of filling or dropping missing values. Filling missing values with any central tendency measure risks losing originality, and dropping them results in losing almost half of the data. For other variables like skin, blood pressure (bp), and body mass index (bmi), excluding glucose (glu), filling missing values with the median is suitable due to the presence of outliers. For glu, filling missing values with the mean is a reasonable approach.

```
[19]: data.drop(['insulin'],axis=1, inplace=True)
      data.shape
```

```
[19]: (768, 8)
```

```
[20]: # Drop missing values for a specific feature

      data.dropna(subset=['skin'], inplace=True)
      sns.heatmap(data.isnull(), cbar=False)
      plt.show()
```



```
[21]: data['glu'].fillna(data["glu"].mean() , inplace = True)
      data['bp'].fillna(data["bp"].median() , inplace = True)
```

```
data['bmi'].fillna(data["bmi"].median() , inplace = True)
```

```
[22]: Total = data.isnull().sum().sort_values(ascending = False)
Percentage = (data.isnull().sum()*100/data.isnull().count()).
↳sort_values(ascending = False)
Missing_Values = pd.concat([Total, Percentage], axis = 1, keys = ['Total',
↳'Percentage of missing observations'])
Missing_Values
```

```
[22]:
```

	Total	Percentage of missing observations
npreg	0	0.0
ped	0	0.0
age	0	0.0
type	0	0.0
glu	0	0.0
bp	0	0.0
skin	0	0.0
bmi	0	0.0

```
[23]: data.shape
```

```
[23]: (541, 8)
```

## 5 4.1 Univariate Analysis

1.glu-Plasma glucose concentration in an oral glucose tolerance test

```
[24]: data.glu.describe()
```

```
[24]: count    541.000000
mean     120.940299
std       30.787626
min       56.000000
25%       99.000000
50%      116.000000
75%      140.000000
max      199.000000
Name: glu, dtype: float64
```

The average plasma glucose concentration in an oral glucose tolerance test is 120.89 mg/dL, ranging from a minimum of 56 mg/dL to a maximum of 199 mg/dL. Among the observations, 25% show a glucose level less than or equal to 99 mg/dL, and 50% have a level less than or equal to 117 mg/dL. A 2-hour plasma glucose level below 140 mg/dL is considered normal. The range of 140-199 mg/dL indicates impaired glucose tolerance, while a level of 200 mg/dL or higher indicates diabetes.

## 6 Skewness and Kurtosis

```
[25]: print("Skewness: %f" % data['glu'].skew())  
      print("Kurtosis: %f" % data['glu'].kurt())
```

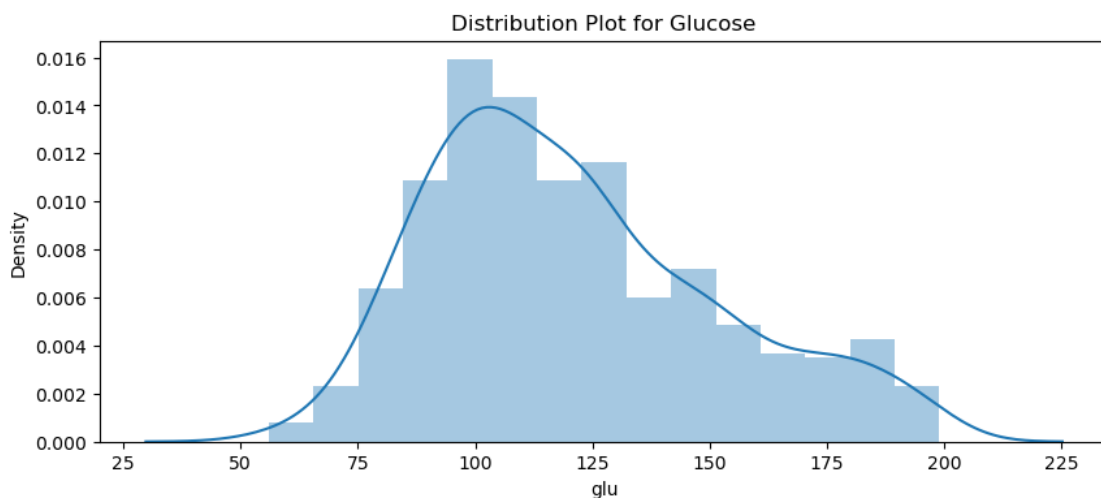
Skewness: 0.617323

Kurtosis: -0.294508


In summary, the skewness value suggests a minor asymmetry to the right in the distribution of plasma glucose concentration, while the kurtosis value indicates that the tails of the distribution are relatively close to those of a normal distribution. These results suggest that the distribution may be roughly symmetric with moderate tails, but a more detailed examination and statistical tests may be needed for a comprehensive assessment of normality.

```
[26]: sns.distplot(data.glu)  
      plt.title("Distribution Plot for Glucose")
```

```
[26]: Text(0.5, 1.0, 'Distribution Plot for Glucose')
```



```
[27]: # q-q plot:q-q plot is used to compare the quantiles of two distributions  
      # p-p plot:p-p plot is the way to visual comparison of cdf of the two  
      ↪distributions  
      import scipy.stats as stats  
      plt.figure(figsize = (8,5))  
      stats.probplot(data["glu"],plot=plt)  
      plt.title("Q-Q Plot for Glucose")  
      plt.show()
```



```
[28]: import numpy as np
from scipy.stats import jarque_bera

# Perform Jarque-Bera test
statistic, p_value = jarque_bera(data.glu)

# Display the results
print(f"Jarque-Bera statistic: {statistic}")
print(f"P-value: {p_value}")

# Check the null hypothesis
if p_value < 0.05:
    print("The glu does not come from a normal distribution (reject the null_
    ↪hypothesis).")
else:
    print("The glu comes from a normal distribution (fail to reject the null_
    ↪hypothesis).")
```

Jarque-Bera statistic: 36.23866327547381

P-value: 1.351681139466535e-08

The glu does not come from a normal distribution (reject the null hypothesis).

The confirmation of non-normal distribution for plasma glucose concentration in an oral glucose

tolerance test is supported by the density plot, Q-Q plot, and Jarque-Bera test.

## 7 2.BP

```
[29]: data.bp.describe()
```

```
[29]: count      541.000000  
      mean       71.463956  
      std        12.261997  
      min        24.000000  
      25%        64.000000  
      50%        72.000000  
      75%        80.000000  
      max        110.000000  
      Name: bp, dtype: float64
```

The average blood pressure is approximately 71.46, reflecting a central tendency, but there is notable variability in the dataset. Quartile values offer insights into blood pressure distribution across different percentiles. The minimum value of 24 suggests potential risks, as extremely low blood pressure is associated with conditions such as severe hypotension, stiff arteries in the elderly, diabetes, arteriovenous malformation, and aortic dissection. The mean value of 71.46 falls within the normal range for blood pressure. Monitoring and understanding such variations are crucial for assessing potential health risks and conditions within the population.

```
[30]: print("Skewness: %f" % data['bp'].skew())  
      print("Kurtosis: %f" % data['bp'].kurt())
```


```
Skewness: -0.007283  
Kurtosis: 0.823193
```

The skewness near zero suggests a nearly symmetric distribution of blood pressure, and the mesokurtic kurtosis indicates a moderate peak and tails, similar to what would be expected in a normal distribution. These characteristics provide insights into the shape and symmetry of the blood pressure distribution in your dataset.


```
[31]: sns.distplot(data.bp)  
      plt.title("Distribution Plot for BP")
```

```
[31]: Text(0.5, 1.0, 'Distribution Plot for BP')
```





```
[32]: plt.figure(figsize = (8,5))  
stats.probplot(data["bp"],plot=plt)  
plt.title("Q-Q Plot for BP")  
plt.show()
```



```
[33]: # Perform Jarque-Bera test
statistic, p_value = jarque_bera(data.bp)

# Display the results
print(f"Jarque-Bera statistic: {statistic}")
print(f"P-value: {p_value}")

# Check the null hypothesis
if p_value < 0.05:
    print("The BP does not come from a normal distribution (reject the null_
    ↪hypothesis).")
else:
    print("The BP comes from a normal distribution (fail to reject the null_
    ↪hypothesis).")
```

Jarque-Bera statistic: 14.595424282189091

P-value: 0.0006770860819314788

The BP does not come from a normal distribution (reject the null hypothesis).

The confirmation of non-normal distribution for Blood Pressure is supported by the density plot, Q-Q plot, and Jarque-Bera test.

## 8 3.Skin Type

```
[34]: data.skin.describe()
```

```
[34]: count      541.000000
mean         29.153420
std          10.476982
min           7.000000
25%          22.000000
50%          29.000000
75%          36.000000
max          99.000000
Name: skin, dtype: float64
```

The “skin thickness” values range from 7 to 99, with an average around 29.15. The data shows some variability, with a moderate spread of values.

```
[35]: print("Skewness: %f" % data['skin'].skew())
print("Kurtosis: %f" % data['skin'].kurt())
```

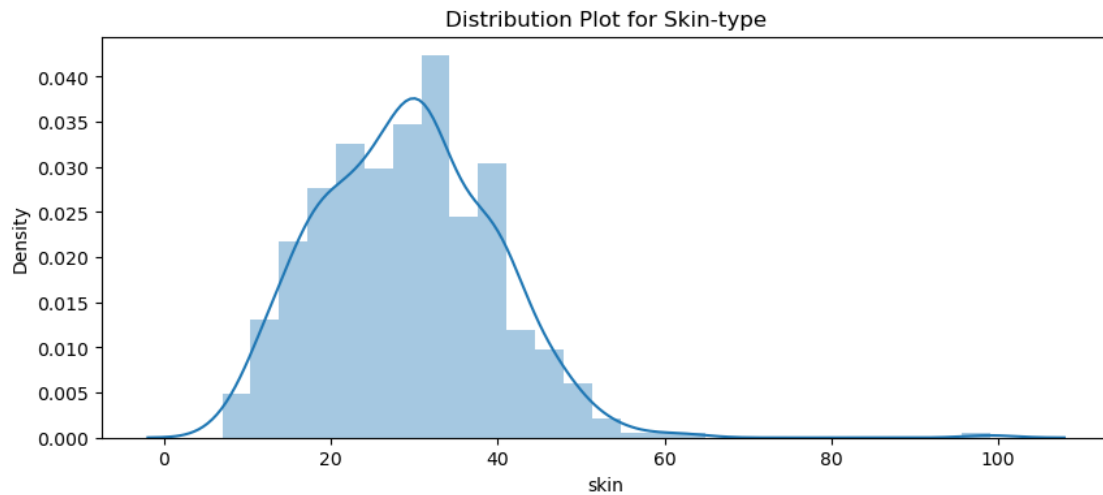
Skewness: 0.690619

Kurtosis: 2.935491

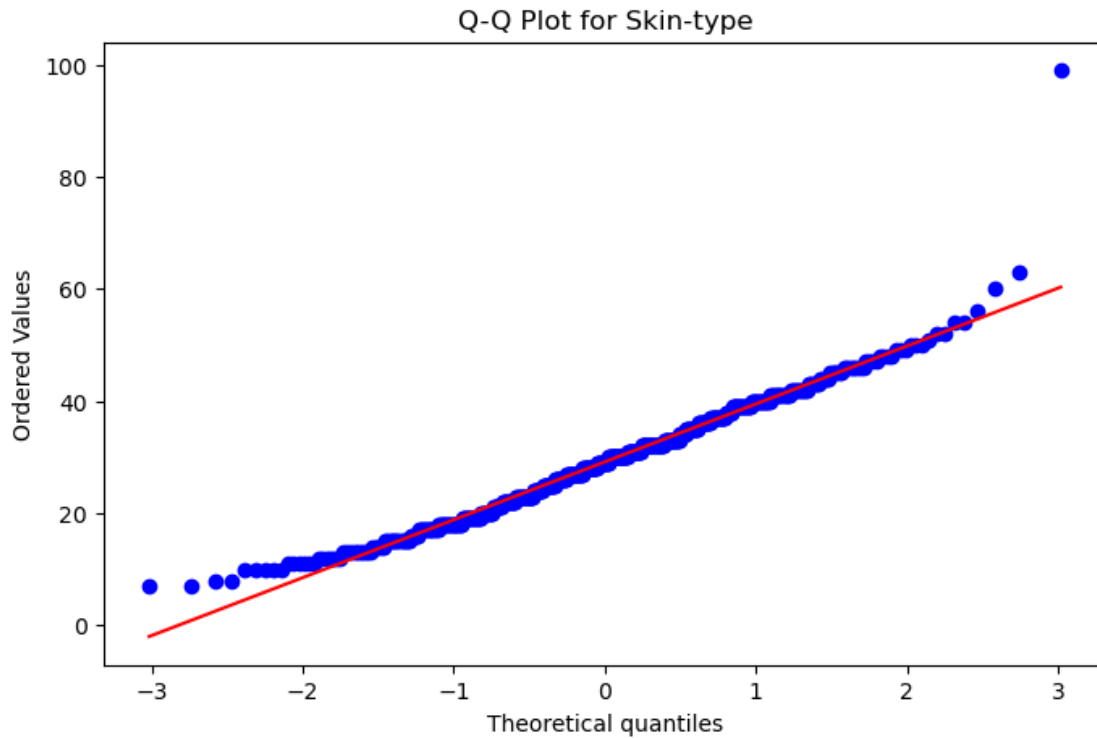
The positive skewness suggests a right-skewed distribution of skin thickness, and the leptokurtic kurtosis indicates heavier tails and a sharper peak. These characteristics provide insights into the asymmetry and concentration of values in the skin thickness distribution in your dataset.

```
[36]: sns.distplot(data.skin)
plt.title("Distribution Plot for Skin-type")
```

```
[36]: Text(0.5, 1.0, 'Distribution Plot for Skin-type')
```



```
[37]: plt.figure(figsize = (8,5))
stats.probplot(data["skin"],plot=plt)
plt.title("Q-Q Plot for Skin-type")
plt.show()
```



```
[38]: # Perform Jarque-Bera test
statistic, p_value = jarque_bera(data.skin)

# Display the results
print(f"Jarque-Bera statistic: {statistic}")
print(f"P-value: {p_value}")

# Check the null hypothesis
if p_value < 0.05:
    print("The Skin-type does not come from a normal distribution (reject the_
    ↪null hypothesis).")
else:
    print("The Skin-type comes from a normal distribution (fail to reject the_
    ↪null hypothesis).")
```

Jarque-Bera statistic: 231.99762905260965

P-value: 4.191359802100761e-51

The Skin-type does not come from a normal distribution (reject the null hypothesis).

The confirmation of non-normal distribution for Skin-type is supported by the density plot, Q-Q plot, and Jarque-Bera test.

## 9 4.Age

```
[39]: data.age.describe()
```

```
[39]: count      541.000000  
      mean       31.558226  
      std       10.743768  
      min       21.000000  
      25%       23.000000  
      50%       28.000000  
      75%       38.000000  
      max       81.000000  
      Name: age, dtype: float64
```

The average age in the study is 31.50 years, ranging from a minimum of 21 years to a maximum of 81 years. Twenty-five percent of women are younger than 23 years, and 50% of women are 28 years old.


```
[40]: print("Skewness: %f" % data['age'].skew())  
      print("Kurtosis: %f" % data['age'].kurt())
```

```
Skewness: 1.269905  
Kurtosis: 1.180848
```


Based on these values, the age distribution in our study is likely to be somewhat skewed to the right with a few older individuals having ages considerably higher than the mean. Additionally, the distribution may exhibit heavier tails and a slightly more peaked shape compared to a normal distribution. Further analysis and visualization of the age distribution, such as a histogram or a density plot, could provide a clearer picture of the data distribution.

```
[41]: sns.distplot(data.age)  
      plt.title("Distribution Plot for Age")
```

```
[41]: Text(0.5, 1.0, 'Distribution Plot for Age')
```



```
[42]: plt.figure(figsize = (8,5))
stats.probplot(data["age"],plot=plt)
plt.title("Q-Q Plot for Age")
plt.show()
```



```
[43]: # Perform Jarque-Bera test
statistic, p_value = jarque_bera(data.age)

# Display the results
print(f"Jarque-Bera statistic: {statistic}")
print(f"P-value: {p_value}")

# Check the null hypothesis
if p_value < 0.05:
    print("The Age does not come from a normal distribution (reject the null_
    ↪hypothesis).")
else:
    print("The Age comes from a normal distribution (fail to reject the null_
    ↪hypothesis).")
```

Jarque-Bera statistic: 174.87643785709037

P-value: 1.061852055024055e-38

The Age does not come from a normal distribution (reject the null hypothesis).

The confirmation of non-normal distribution for age is supported by the density plot, Q-Q plot, and Jarque-Bera test.

## 10 5.BMI

```
[44]: data.bmi.describe()
```

```
[44]: count      541.000000
mean        32.895379
std         6.859116
min         18.200000
25%         27.900000
50%         32.800000
75%         36.900000
max         67.100000
Name: bmi, dtype: float64
```

```
[45]: # Calculate mode
mode_result = stats.mode(data.bmi)

print("Mode:", mode_result.mode[0])
```

Mode: 32.0

The mean Body Mass Index (BMI), computed as the ratio of weight in kilograms to the square of height in meters, is 32.89. The BMI ranges from a minimum of 18.20 to a maximum of 67.10. Notably, approximately 50% of women exhibit a BMI below 32.80, a value that aligns closely with both the mean and mode of the BMI. This indicates a close similarity between the mean, median, and mode, suggesting a distribution with a central tendency.

```
[46]: print("Skewness: %f" % data['bmi'].skew())
      print("Kurtosis: %f" % data['bmi'].kurt())
```


Skewness: 0.626830

Kurtosis: 1.272595

The distribution is right-skewed, suggesting a possible concentration of individuals with higher BMI values. The kurtosis indicates that the distribution has heavier tails and is more peaked than a normal distribution, possibly indicating the presence of more extreme BMI values.

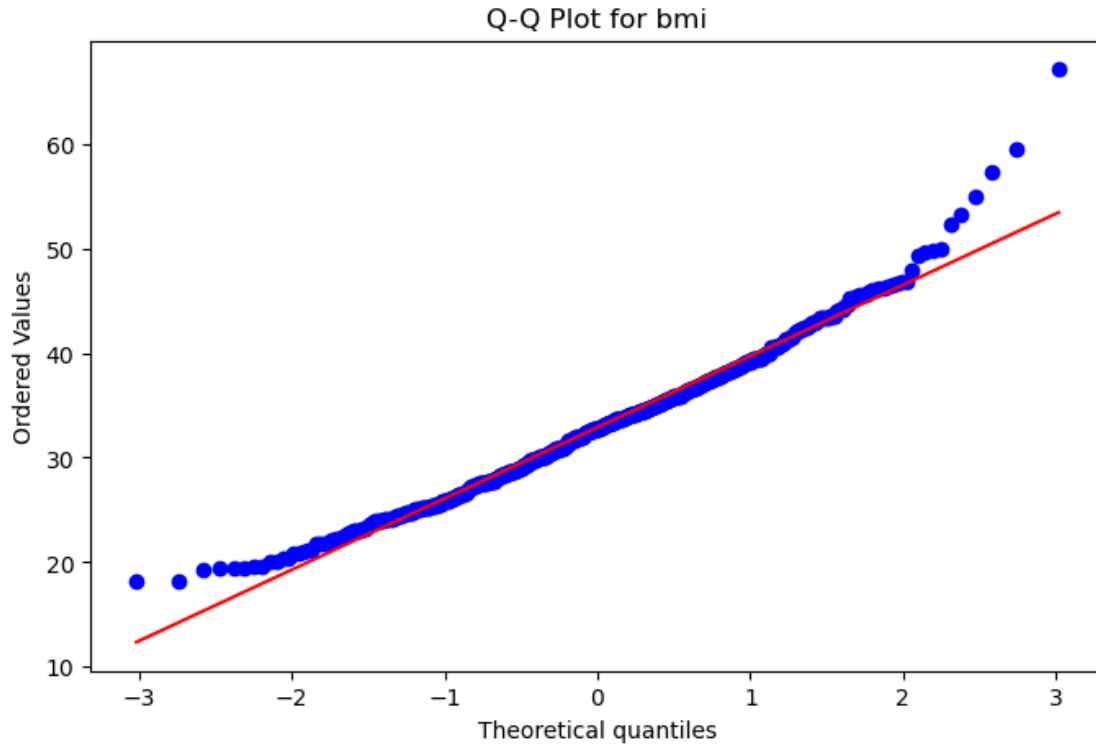
```
[47]: sns.distplot(data.bmi)
      plt.title("Distribution Plot for BMI")
```

```
[47]: Text(0.5, 1.0, 'Distribution Plot for BMI')
```



```
[48]: plt.figure(figsize = (8,5))
      stats.probplot(data["bmi"],plot=plt)
      plt.title("Q-Q Plot for bmi")
      plt.show()
```





```
[49]: # Perform Jarque-Bera test
statistic, p_value = jarque_bera(data.bmi)

# Display the results
print(f"Jarque-Bera statistic: {statistic}")
print(f"P-value: {p_value}")

# Check the null hypothesis
if p_value < 0.05:
    print("The BMI does not come from a normal distribution (reject the null_
    ↪hypothesis).")
else:
    print("The BMI comes from a normal distribution (fail to reject the null_
    ↪hypothesis).")
```

Jarque-Bera statistic: 70.44127573826687

P-value: 5.056748150849164e-16

The BMI does not come from a normal distribution (reject the null hypothesis).

The confirmation of non-normal distribution for BMI is supported by the density plot, Q-Q plot, and Jarque-Bera test.

## 11 6. Number of Pregnancies

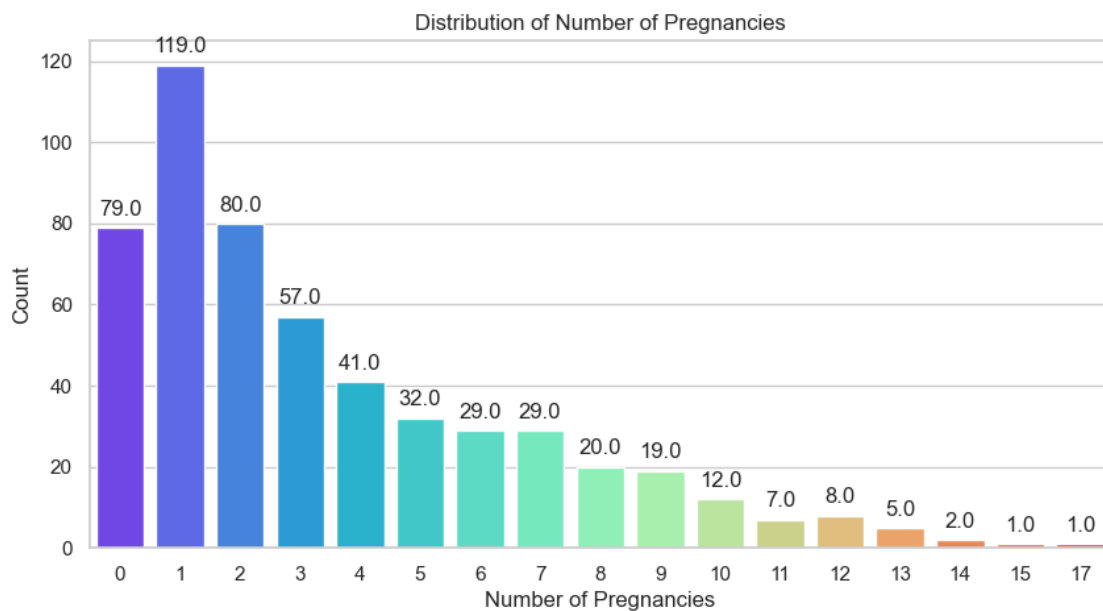
```
[50]: sns.set(style="whitegrid")
plt.figure(figsize=(10, 5))

# Create the countplot
ax = sns.countplot(x=data.npreg, palette="rainbow")

# Add legends, labels, and values on each bar
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
    ↪get_height()),
    ↪ha='center', va='center', xytext=(0, 10), textcoords='offset_
    ↪points')

plt.xlabel("Number of Pregnancies")
plt.ylabel("Count")
plt.title("Distribution of Number of Pregnancies")

plt.show()
```




Where 79 women have no children and 119 women have one child, the mode of the distribution is determined to be 1.

## 12 7. Type (outcome): Diabetic (Yes or No)

```
[51]: #Splitting the DataFrame into feature variables (data_x) and the target_
      ↪variable (data_y).
data_x = data.iloc[:, data.columns != 'type']
data_y = data.iloc[:,data.columns == 'type']
```

```
[52]: class_frequency =data_y.type.value_counts()
sns.countplot(data=data_y,x ="type")
plt.text(x = -0.05, y =data_y.value_counts()[0]+1, s =_
      ↪str(round((class_frequency[0])*100/len(data_y),2)) + '%')
plt.text(x = 0.95, y =data_y.value_counts()[1]+1, s =_
      ↪str(round((class_frequency[1])*100/len(data_y),2)) + '%')
plt.title('Count Plot for Target Variable Type', fontsize = 15)
plt.xlabel('Target Variable (Type)', fontsize = 15)
plt.ylabel('Count', fontsize = 15)
plt.show()
```



Conclusion: In our research study, we have a target variable called “type,” which serves as an indicator of an individual’s diabetes status. This variable takes on two distinct values:

When “type” is assigned the value 0, it signifies that the individual is categorized as non-diabetic. In other words, individuals with this value do not have diabetes.

When “type” is assigned the value 1, it indicates that the individual is categorized as diabetic. This means they have been diagnosed with diabetes.

Based on our analysis of the data, we have found that 33.27% of the study participants fall into the diabetic category (type = 1), indicating that they have diabetes. In contrast, 66.73% of the participants fall into the non-diabetic category (type = 0), signifying that they do not have diabetes.

These percentages are important findings as they provide insights into the prevalence of diabetes within our study population. It suggests that a significant portion of the individuals in our study have been diagnosed with diabetes (33.27%), while the majority are non-diabetic (66.73%). This information is crucial for understanding the distribution of diabetes within the studied group and may have implications for healthcare interventions or further research related to diabetes management and prevention.

## 13 4.2 Multivariate Analysis

### 14 1.Box Plots for Target Variable (Type) with Different Features

```
[53]: import seaborn as sns
import matplotlib.pyplot as plt

# Assuming 'data' is your DataFrame containing the variables
# 'type', 'bmi', 'age', 'npreg', 'skin', 'glu', 'bp'

# Set up the figure with subplots
fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(15, 15))

# Boxplot for 'type' vs 'bmi'
sns.boxplot(x='type', y='bmi', data=data, ax=axes[0, 0])
axes[0, 0].set_title('Boxplot: Type vs BMI')

# Boxplot for 'type' vs 'age'
sns.boxplot(x='type', y='age', data=data, ax=axes[0, 1])
axes[0, 1].set_title('Boxplot: Type vs Age')

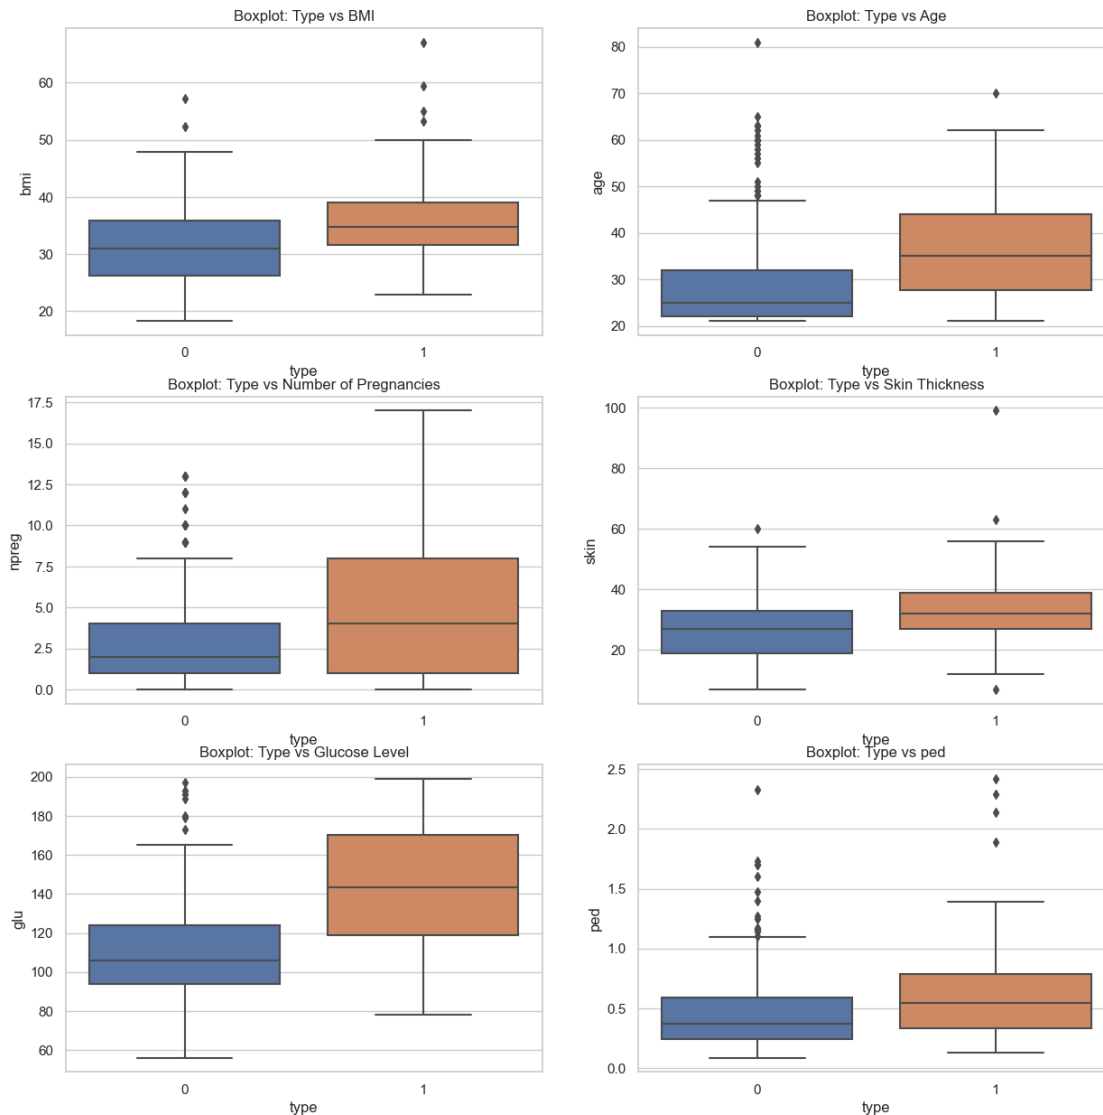
# Boxplot for 'type' vs 'npreg'
sns.boxplot(x='type', y='npreg', data=data, ax=axes[1, 0])
axes[1, 0].set_title('Boxplot: Type vs Number of Pregnancies')

# Boxplot for 'type' vs 'skin'
sns.boxplot(x='type', y='skin', data=data, ax=axes[1, 1])
axes[1, 1].set_title('Boxplot: Type vs Skin Thickness')

# Boxplot for 'type' vs 'glu'
sns.boxplot(x='type', y='glu', data=data, ax=axes[2, 0])
axes[2, 0].set_title('Boxplot: Type vs Glucose Level')

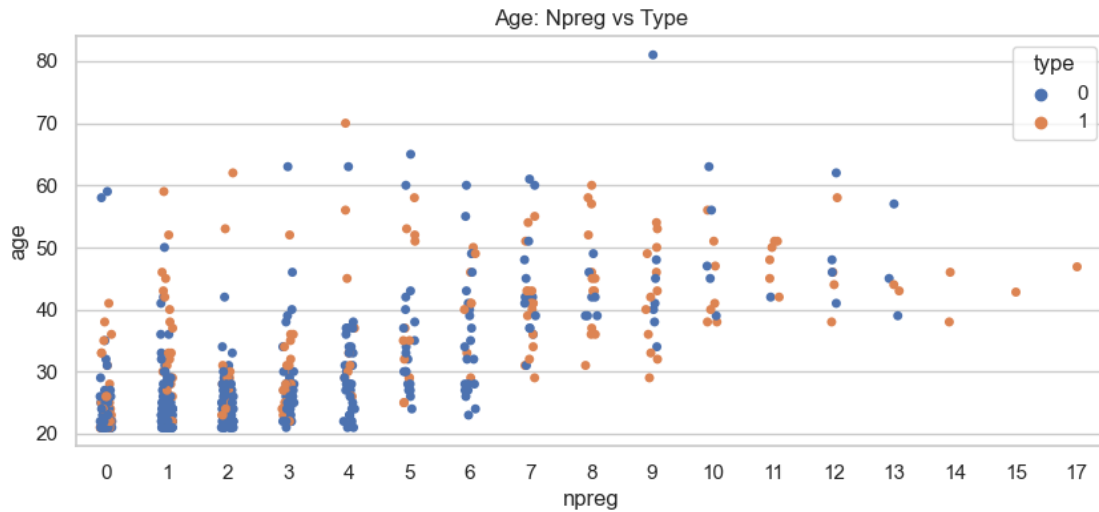
# Boxplot for 'type' vs 'ped'
sns.boxplot(x='type', y='ped', data=data, ax=axes[2, 1])
axes[2, 1].set_title('Boxplot: Type vs ped')

# For example, if using matplotlib
plt.savefig('my_plot.png', bbox_inches='tight')
```



From the multiple boxplots presented above, a notable observation emerges: individuals with diabetes exhibit significantly higher mean values for BMI, Glucose, Pedigree Function (ped), and age compared to those without diabetes. Specifically, the boxplots reveal a clear distinction in these variables between the two groups. The higher mean values in BMI, Glucose, ped, and age for individuals with diabetes suggest potential associations or characteristics that differentiate the diabetic and non-diabetic populations in the dataset.

```
[54]: sns.stripplot(x="npreg",y="age",data=data,hue="type")
plt.title('Age: Npreg vs Type')
plt.show()
```




The likelihood of diabetes appears to be higher in cases of a greater number of pregnancies than in instances with fewer pregnancies. This observation implies a potential correlation between the number of pregnancies and the risk of diabetes among women. Further analysis and statistical testing would be needed to establish a conclusive relationship and determine potential contributing factors.

## 15 2. Analysis of ped with glucose

```
[55]: x = data[['ped', 'glu']]
      sns.jointplot(x=x.loc[:, 'glu'], y=x.loc[:, 'ped'], kind="reg", color="#ce1414")
```

```
[55]: <seaborn.axisgrid.JointGrid at 0x19f20f38c70>
```



```
[56]: data_x.corr()
```


```
[56]:
```

	npreg	ped	age	glu	bp	skin	bmi
npreg	1.000000	0.003539	0.644686	0.127319	0.205066	0.100239	0.019495
ped	0.003539	1.000000	0.066834	0.160832	0.006823	0.115016	0.150885
age	0.644686	0.066834	1.000000	0.279718	0.347147	0.166816	0.081580
glu	0.127319	0.160832	0.279718	1.000000	0.217581	0.227369	0.247116
bp	0.205066	0.006823	0.347147	0.217581	1.000000	0.226723	0.309508
skin	0.100239	0.115016	0.166816	0.227369	0.226723	1.000000	0.647828
bmi	0.019495	0.150885	0.081580	0.247116	0.309508	0.647828	1.000000

```
[57]: corr=data_x.corr()
```

```
sns.heatmap(corr, cmap = 'YlGnBu', vmax = 1.0, vmin = -1.0, annot = True,
↪annot_kws = {"size": 12})
```

[57]: <Axes: >



Age and number of pregnancy: There is a positive correlation between Age and number of pregnancy. This means that as those who are having higher age more number of children. BMI and Skin thickness: There is a positive correlation between BMI and Skin thickness. This means that as those who are having more BMI with higher skin thickness.

Remaining Variables: For the other variables in our dataset, there are weak or no significant correlations between them. This implies that changes in one of these variables do not strongly influence or predict changes in the others. They are relatively independent of each other in terms of their relationships within the dataset.

Understanding these correlations is important in various fields, such as medicine or statistics, as it helps us identify potential relationships between variables. In this case, the strong positive correlation between diastolic and systolic blood pressure might indicate that they are both influenced by similar factors or share a common physiological basis. Conversely, the weak or non-existent correlations among the remaining variables suggest that they are relatively unrelated in the context of this dataset. Please correct the sentence

[58]: `data.shape`

[58]: (541, 8)

```
[59]: data_numeric = data.select_dtypes(include=np.number)
print(data_numeric.columns)
data_categoric = data.select_dtypes(include = object)
print(data_categoric.columns)
```



```
Index(['npreg', 'ped', 'age', 'type', 'glu', 'bp', 'skin', 'bmi'],
      dtype='object')
Index([], dtype='object')
```

```
[60]: #dummy_variables = pd.get_dummies(data_categoric, drop_first = True)
```

```
[61]: #data_dummy = pd.concat([data_numeric, dummy_variables], axis=1)
      data_dummy=data
```

```
[62]: data_dummy.head(1)
```

```
[62]:      npreg      ped  age  type      glu      bp  skin      bmi
      0         6  0.627   50      1  148.0   72.0   35.0   33.6
```

```
[63]: X = data_dummy.drop(['type'], axis = 1)
      X=sm.add_constant(X)
      y = pd.DataFrame(data_dummy['type'])
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
      ↪random_state = 1)
```

```
[64]: def get_test_report(model):
      return(classification_report(y_test,y_pred))
```

```
[65]: def kappa_score(model):
      return(cohen_kappa_score(y_test,y_pred))
```

```
[66]: def plot_confusion_matrix(model):
      cm = confusion_matrix(y_test, y_pred)
      conf_matrix = pd.DataFrame(data = cm,columns = ['Predicted:0', 'Predicted:
      ↪1'], index = ['Actual:0', 'Actual:1'])
      sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap =
      ↪ListedColormap(['lightskyblue']),cbar = False, linewidths = 0.1, annot_kws =
      ↪{'size':25})
      plt.xticks(fontsize = 20)
      plt.yticks(fontsize = 20)
      plt.show()
```

```
[67]: def plot_roc(model):
      fpr,tpr,_=roc_curve(y_test,y_pred_prob)
      plt.plot(fpr,tpr)
      plt.xlim([0.0,1.0])
      plt.ylim([0.0,1.0])
      plt.plot([0,1],[0,1],"r--")
      plt.title("ROC Curve",fontsize=15)
      plt.xlabel("False positive",fontsize=15)
      plt.ylabel("True positive",fontsize=15)
```

```
plt.text(x=0.02,y=0.9,s=("AUC Score:
↪",round(roc_auc_score(y_test,y_pred_prob),4)))
plt.grid(True)
```

```
[68]: score_card=pd.DataFrame(columns=["Model","AUC Score","Precision Score","Recall_
↪Score","Accuracy Score","Kappa Score","f1-Score"])
def update_score_card(model_name):
    global score_card
    score_card=score_card.append({"Model":model_name,"AUC Score":
↪roc_auc_score(y_test,y_pred_prob),"Precision Score":metrics.
↪precision_score(y_test,y_pred),"Recall Score":metrics.
↪accuracy_score(y_test,y_pred),'Accuracy Score': metrics.
↪accuracy_score(y_test, y_pred),"Kappa Score":
↪cohen_kappa_score(y_test,y_pred),"f1-Score":metrics.
↪f1_score(y_test,y_pred)},ignore_index=True)
    return(score_card)
```

After completing data cleaning and certain exploratory data analysis (EDA) steps, we partitioned the data into two sets: a training set comprising 80% of the observations and a test set with 20% of the observations to assess the model's accuracy.

In this phase, we applied various machine learning models, namely Logistic Regression, Decision Tree, Naive Bayes, and Support Vector Machine. Subsequently, we compared the accuracy of these different models, selecting the best-performing ones for deployment.

```
[69]: Log_Reg_Full_Model=sm.Logit(y_train,X_train).fit()
print(Log_Reg_Full_Model.summary())
```

Optimization terminated successfully.

Current function value: 0.436427

Iterations 7

#### Logit Regression Results

```
=====
Dep. Variable:          type    No. Observations:          432
Model:                  Logit    Df Residuals:            424
Method:                  MLE     Df Model:              7
Date:                   Sun, 28 Jan 2024    Pseudo R-squ.:        0.3090
Time:                   16:32:22    Log-Likelihood:       -188.54
converged:              True    LL-Null:             -272.85
Covariance Type:        nonrobust    LLR p-value:         4.899e-33
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-9.7352	1.123	-8.667	0.000	-11.937	-7.534
npreg	0.1335	0.047	2.824	0.005	0.041	0.226
ped	1.3995	0.401	3.491	0.000	0.614	2.185
age	0.0261	0.015	1.692	0.091	-0.004	0.056
glu	0.0369	0.005	7.616	0.000	0.027	0.046

bp	-0.0126	0.012	-1.088	0.277	-0.035	0.010
skin	0.0034	0.017	0.204	0.838	-0.030	0.036
bmi	0.0905	0.026	3.491	0.000	0.040	0.141
=====						

Except for BP and skin thickness, all variables are deemed statistically significant, as their p-values exceed 0.05 for intercept, npreg, ped, age, glu, and bmi. This implies significance at the 5% level. However, the inclusion of non-significant variables raises concerns about potential overfitting. To mitigate this, we intend to select only significant variables using forward or backward elimination in feature selection methods and subsequently assess the model's accuracy.

```
[70]: y_pred_prob=Log_Reg_Full_Model.predict(X_test)
y_pred=["0" if x<0.5 else "1" for x in y_pred_prob]
y_pred=np.array(y_pred,dtype=np.float32)
y_pred[0:5]
plot_confusion_matrix(Log_Reg_Full_Model)
plot_roc(Log_Reg_Full_Model)
update_score_card(model_name="Logistic_Regression with Full Model")
```


Actual:	Actual:0	66	4
	Actual:1	19	20
		Predicted:0	Predicted:1

```
[70]:
```

	Model	AUC Score	Precision Score	\
0	Logistic_Regression with Full Model	0.864835	0.833333	

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921



The confusion matrix indicates a 17.43% false negative and a 3.6% false positive, resulting in a total accuracy of 78.89%.

the AUC score of 0.8648 indicates that there is a high probability that the model will assign a higher predicted probability to a randomly chosen positive instance compared to a randomly chosen negative instance. The closer the AUC score is to 1, the better the model's ability to distinguish between positive and negative instances.

```
[71]: #SGDC Classifier with constant(intercept term alpha)
SGD = SGDClassifier(loss = 'log', random_state = 10)
Log_Reg_with_SGD = SGD.fit(X_train, y_train)
```

```
[72]: y_pred_prob =Log_Reg_with_SGD.predict_proba(X_test)[: ,1]
      y_pred_prob
```


```
[72]: array([0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1., 1., 0., 1., 0., 0.,
            0., 0., 1., 0., 0., 1., 1., 1., 1., 0., 0., 0., 0., 0., 0., 0., 0.,
            0., 0., 0., 1., 1., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1.,
            0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 1.,
            0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0.,
            0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0., 1.,
            0., 0., 0., 1., 0., 0., 0.] )
```

```
[73]: y_pred =Log_Reg_with_SGD.predict(X_test)
      y_pred
```

```
[73]: array([0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0,  
          1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,  
          0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,  
          0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0])
```

```
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0],
dtype=int64)
```

```
[74]: plot_confusion_matrix(Log_Reg_with_SGD)
```



The confusion matrix reveals a 22.93% false negative rate and a 7.3% false positive rate, leading to an overall accuracy of 69.72%. This accuracy is comparatively lower than that of the previous model.


```
[75]: test_report = get_test_report(Log_Reg_with_SGD)
print(test_report)
```

	precision	recall	f1-score	support
0	0.71	0.89	0.79	70
1	0.64	0.36	0.46	39
accuracy			0.70	109
macro avg	0.67	0.62	0.62	109
weighted avg	0.69	0.70	0.67	109

```
[76]: kappa_value = kappa_score(Log_Reg_with_SGD)
print(kappa_value)
```

```
0.2708291100750051
```

```
[77]: plot_roc(Log_Reg_with_SGD)
```



An Area Under the Curve (AUC) score of 0.6427 on the Receiver Operating Characteristic (ROC) curve suggests a moderate discriminatory performance of the model. The ROC curve illustrates the trade-off between the true positive rate (sensitivity) and the false positive rate (1-specificity) across various threshold values.

```
[78]: update_score_card(model_name = 'Logistic Regression (SGD)')
```

```
[78]:
```

	Model	AUC Score	Precision Score \
0	Logistic Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016

```
[79]: X = data_dummy.drop(['type'], axis = 1)
y = pd.DataFrame(data_dummy['type'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
random_state = 1)
```

```
[80]: Log_Reg_Without_Intercept=sm.Logit(y_train,X_train).fit()
print(Log_Reg_Without_Intercept.summary())
```

Optimization terminated successfully.

Current function value: 0.560074

Iterations 6

#### Logit Regression Results

```
=====
Dep. Variable:          type    No. Observations:          432
Model:                Logit    Df Residuals:              425
```

Method:	MLE	Df Model:	6
Date:	Sun, 28 Jan 2024	Pseudo R-squ.:	0.1132
Time:	16:32:28	Log-Likelihood:	-241.95
converged:	True	LL-Null:	-272.85
Covariance Type:	nonrobust	LLR p-value:	1.946e-11

	coef	std err	z	P> z	[0.025	0.975]
npreg	0.1354	0.043	3.168	0.002	0.052	0.219
ped	0.6483	0.341	1.899	0.058	-0.021	1.317
age	0.0054	0.014	0.378	0.705	-0.023	0.033
glu	0.0206	0.004	5.159	0.000	0.013	0.028
bp	-0.0593	0.010	-6.229	0.000	-0.078	-0.041
skin	0.0152	0.015	1.004	0.315	-0.014	0.045
bmi	-0.0114	0.022	-0.528	0.598	-0.054	0.031


```
[81]: y_pred_prob=Log_Reg_Without_Intercept.predict(X_test)
y_pred=["0" if x<0.5 else "1" for x in y_pred_prob]
y_pred=np.array(y_pred,dtype=np.float32)
y_pred[0:5]
plot_confusion_matrix(Log_Reg_Without_Intercept)
plot_roc(Log_Reg_Without_Intercept)
update_score_card(model_name="Log_Reg_Without_Intercept")
```

Actual:0	59	11
	22	17
	Predicted:0	Predicted:1

	Model	AUC Score	Precision Score \
0	Logistic Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143

Recall Score	Accuracy Score	Kappa Score	f1-Score
--------------	----------------	-------------	----------

0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463



The confusion matrix reveals a 20.18% false negative rate and a 10.09% false positive rate, resulting in an overall accuracy of 69.72%. This accuracy is comparatively lower than that of the first model, suggesting a potential need to consider adding the intercept term. An AUC score of 0.74 implies a fair discriminatory performance, and further analysis or adjustments may be considered to potentially enhance the model's effectiveness.

## 16 Backward Model Selection Using Univariate Statistical Testing

```
[82]: import statsmodels.api as sm
import pandas as pd

# Assume 'df' is your DataFrame with the target variable ('y') and predictor
# variables

# Backward elimination function
def backward_elimination(data, target):
    features = list(data.columns)
    features.remove(target)

    while len(features) > 0:
        model = sm.Logit(data[target], sm.add_constant(data[features]))
        result = model.fit(dis= False)
        max_pvalue = result.pvalues.idxmax()
```



```

        # If the highest p-value is greater than a threshold (e.g., 0.05),
        ↪ remove the corresponding feature
        if result.pvalues[max_pvalue] > 0.05:
            features.remove(max_pvalue)
        else:
            break # If all p-values are below the threshold, stop

    return features

# Example usage
target_variable = 'type'
selected_features_backward = backward_elimination(data_dummy, target_variable)

print("Selected Features (Backward):", selected_features_backward)

```

Selected Features (Backward): ['npreg', 'ped', 'glu', 'bmi']

```

[83]: X = data_dummy.drop(['type', "skin", "age", "bp"], axis = 1)
      X=sm.add_constant(X)
      y = pd.DataFrame(data_dummy['type'])
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
      ↪ random_state = 1)

```

```

[84]: Log_Reg_Backward_Model_Selection=sm.Logit(y_train,X_train).fit()
      print(Log_Reg_Backward_Model_Selection.summary())

```

Optimization terminated successfully.

Current function value: 0.440269

Iterations 6

#### Logit Regression Results

```

=====
Dep. Variable:                type    No. Observations:                432
Model:                        Logit    Df Residuals:                  427
Method:                        MLE     Df Model:                      4
Date:                         Sun, 28 Jan 2024    Pseudo R-squ.:                0.3029
Time:                         16:32:34    Log-Likelihood:               -190.20
converged:                     True    LL-Null:                      -272.85
Covariance Type:               nonrobust    LLR p-value:                  1.065e-34
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	-9.8095	1.004	-9.766	0.000	-11.778	-7.841
npreg	0.1795	0.037	4.836	0.000	0.107	0.252
ped	1.4627	0.395	3.705	0.000	0.689	2.236
glu	0.0377	0.005	8.004	0.000	0.028	0.047
bmi	0.0856	0.020	4.293	0.000	0.047	0.125

```

=====

```

In the model summary table, all variables are found to be significant at any given level of significance. Models of this type are less prone to overfitting. If the model performs well under the performance metrics, we can consider finalizing it for real-world predictions of diabetes

```
[85]: y_pred_prob=Log_Reg_Backward_Model_Selection.predict(X_test)
y_pred=["0" if x<0.5 else "1" for x in y_pred_prob]
y_pred=np.array(y_pred,dtype=np.float32)
y_pred[0:5]
plot_confusion_matrix(Log_Reg_Backward_Model_Selection)
plot_roc(Log_Reg_Backward_Model_Selection)
update_score_card(model_name="Log_Reg_Backward_Model_Selection")
```


Actual:	Actual:0	66	4
	Actual:1	19	20
		Predicted:0	Predicted:1

```
[85]:
```

	Model	AUC Score	Precision Score	\
0	Logistic_Regression with Full Model	0.864835	0.833333	
1	Logistic Regression (SGD)	0.642674	0.636364	
2	Log_Reg_Without_Intercept	0.744689	0.607143	
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333	

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921



The confusion matrix reveals a 17.43% false negative rate and a 3.6% false positive rate, leading to an overall accuracy of 78.89%. The performance of this model is comparable to that of the first model. An AUC score of 0.8648 suggests that the model has a strong ability to discriminate between classes, making it a promising indicator of its overall performance.

```
[86]: X = data_dummy.drop(['type', "skin", "bp", "npreg"], axis = 1)
X=sm.add_constant(X)
y = pd.DataFrame(data_dummy['type'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
random_state = 1)
Log_Reg_With_Domain_Knowledge=sm.Logit(y_train,X_train).fit()
print(Log_Reg_With_Domain_Knowledge.summary())
```

Optimization terminated successfully.

Current function value: 0.448043

Iterations 6

#### Logit Regression Results

=====						
Dep. Variable:		type	No. Observations:		432	
Model:		Logit	Df Residuals:		427	
Method:		MLE	Df Model:		4	
Date:		Sun, 28 Jan 2024	Pseudo R-squ.:		0.2906	
Time:		16:32:37	Log-Likelihood:		-193.55	
converged:		True	LL-Null:		-272.85	
Covariance Type:		nonrobust	LLR p-value:		2.939e-33	
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
const	-10.1884	1.033	-9.866	0.000	-12.212	-8.164
ped	1.2914	0.389	3.316	0.001	0.528	2.055

age	0.0483	0.012	4.198	0.000	0.026	0.071
glu	0.0348	0.005	7.421	0.000	0.026	0.044
bmi	0.0832	0.020	4.258	0.000	0.045	0.122

=====

```
[87]: y_pred_prob=Log_Reg_With_Domain_Knowledge.predict(X_test)
y_pred=["0" if x<0.5 else "1" for x in y_pred_prob]
y_pred=np.array(y_pred,dtype=np.float32)
#plot_confusion_matrix(Log_Reg_With_Domain_Knowledge)
#plot_roc(Log_Reg_With_Domain_Knowledge)
#update_score_card(model_name="Log_Reg_With_Domain_Knowledge")
```

## 17 Decision Tree Classification

```
[88]: X = data_dummy.drop(['type'], axis = 1)
y = pd.DataFrame(data_dummy['type'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳random_state = 1)
```

```
[89]: tuned_parameters=[{"criterion":["gini","entropy"],"min_samples_split":
↳[10,20,30],"max_depth":[3,5,7,9],"min_samples_leaf":
↳[15,20,25,30,35],"max_leaf_nodes":[5,10,15,20,25]}]
```

```
[90]: decision_tree_classification=DecisionTreeClassifier(random_state=10)
grid=GridSearchCV(estimator=decision_tree_classification,param_grid=tuned_parameters,cv=10)
dt_grid=grid.fit(X_train,y_train)
print("Best parameters for DT:",dt_grid.best_params_,"\n")
```

Best parameters for DT: {'criterion': 'entropy', 'max\_depth': 5,  
'max\_leaf\_nodes': 10, 'min\_samples\_leaf': 15, 'min\_samples\_split': 10}

```
[91]: dt_grid_model=DecisionTreeClassifier(criterion=dt_grid.best_params_.
↳get("criterion"),max_depth=dt_grid.best_params_.
↳get("max_depth"),max_leaf_nodes=dt_grid.best_params_.
↳get("max_leaf_nodes"),min_samples_leaf=dt_grid.best_params_.
↳get("min_samples_leaf"),min_samples_split=dt_grid.best_params_.
↳get("min_samples_split"))
```

```
[92]: decision_tree_grid=dt_grid_model.fit(X_train,y_train)
```

```
[93]: y_pred_prob=decision_tree_grid.predict_proba(X_test)[:,-1]
```

```
[94]: y_pred=decision_tree_grid.predict(X_test)
```

```
[95]: plot_confusion_matrix(decision_tree_grid)
```

Actual:	Actual:0	89	16
	Actual:1	18	40
		Predicted:0	Predicted:1

```
[96]: test_report = get_test_report(decision_tree_grid)

# print the performace measures
print(test_report)
```


	precision	recall	f1-score	support
0	0.83	0.85	0.84	105
1	0.71	0.69	0.70	58
accuracy			0.79	163
macro avg	0.77	0.77	0.77	163
weighted avg	0.79	0.79	0.79	163

```
[97]: kappa_value = kappa_score(decision_tree_grid)

# print the kappa value
print(kappa_value)
```

0.5414529207347345

```
[98]: plot_roc(decision_tree_grid)
```



```
[99]: update_score_card(model_name = 'decision_tree_grid')
```

```
[99]:
```

	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754

```
[100]: X = data_dummy.drop(['type'], axis = 1)
y = pd.DataFrame(data_dummy['type'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳ random_state = 1)
```

```
[101]: from sklearn.naive_bayes import GaussianNB
```


```
[102]: Naive_Bayes_Model =GaussianNB().fit(X_train, y_train)
```

```
[103]: y_pred_prob =Naive_Bayes_Model .predict_proba(X_test)[: ,1]
```

```
[104]: y_pred = Naive_Bayes_Model.predict(X_test)
y_pred[0:11]
```

```
[104]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1], dtype=int64)
```


```
[105]: plot_confusion_matrix(Naive_Bayes_Model)
```



```
[106]: test_report = get_test_report(Naive_Bayes_Model)
print(test_report)
```

	precision	recall	f1-score	support
0	0.81	0.87	0.83	105
1	0.72	0.62	0.67	58
accuracy			0.78	163
macro avg	0.76	0.74	0.75	163
weighted avg	0.77	0.78	0.78	163

```
[107]: plot_roc(Naive_Bayes_Model)
```



```
[108]: update_score_card(model_name = 'Naive_Bayes_Model')
```

```
[108]:
```

	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286
5	Naive_Bayes_Model	0.847291	0.720000

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754
5	0.779141	0.779141	0.502880	0.666667

```
[109]: from sklearn.svm import SVC
```

```
[110]: svc_linear = SVC(kernel='linear', probability=True) # Specify
        ↳ 'probability=True' to enable probability estimates
        svm_linear=svc_linear.fit(X_train, y_train)
        y_pred_prob =svm_linear.predict_proba(X_test)[:,-1]
        y_pred =svm_linear .predict(X_test)
        plot_confusion_matrix(svm_linear)
        test_report = get_test_report(svm_linear)
        print(test_report)
        plot_roc(svm_linear)
```



```
update_score_card(model_name = 'svm_linear')
```


Actual:	Actual:0	95	10
	Actual:1	25	33
		Predicted:0	Predicted:1

	precision	recall	f1-score	support
0	0.79	0.90	0.84	105
1	0.77	0.57	0.65	58
accuracy			0.79	163
macro avg	0.78	0.74	0.75	163
weighted avg	0.78	0.79	0.78	163

```
[110]:
```

	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286
5	Naive_Bayes_Model	0.847291	0.720000
6	svm_linear	0.838095	0.767442

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754
5	0.779141	0.779141	0.502880	0.666667
6	0.785276	0.785276	0.502832	0.653465



```
[111]: svc_poly = SVC(kernel='poly', probability=True) # Specify 'probability=True'
        ↳to enable probability estimates
        svm_poly=svc_poly.fit(X_train, y_train)
        y_pred_prob =svm_poly.predict_proba(X_test)[:,-1]
        y_pred =svm_poly .predict(X_test)
        plot_confusion_matrix(svm_poly)
        test_report = get_test_report(svm_poly)
        print(test_report)
        plot_roc(svm_poly)
        update_score_card(model_name = 'svm_poly')
```

Actual:0	100	5
Actual:1	28	30
	Predicted:0	Predicted:1


precision    recall    f1-score    support

	0	0.78	0.95	0.86	105
	1	0.86	0.52	0.65	58
accuracy				0.80	163
macro avg		0.82	0.73	0.75	163
weighted avg		0.81	0.80	0.78	163

[111]:

	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286
5	Naive_Bayes_Model	0.847291	0.720000
6	svm_linear	0.838095	0.767442
7	svm_poly	0.849261	0.857143

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754
5	0.779141	0.779141	0.502880	0.666667
6	0.785276	0.785276	0.502832	0.653465
7	0.797546	0.797546	0.515362	0.645161



The table above compares the performance of seven models. Among these models, except for the Logistic model with SGD classifier, the remaining models exhibit nearly equal performance. Some

models, such as decision tree, SVM, and naive Bayes, are complex (high variance) and may be prone to overfitting. When simple models (low variance) perform nearly as well as complex models, it is preferable to choose the simpler ones. We are selecting Logistic Regression with Backward Elimination method because all variables are significant, and it involves fewer features. To mitigate the risk of high variance, opting for simpler models is advisable, especially when the model does not exhibit signs of underfitting.

```
[112]: data_dummy.head(2)
```

```
[112]:      npreg    ped  age  type    glu    bp  skin    bmi
0         6  0.627   50     1  148.0  72.0  35.0  33.6
1         1  0.351   31     0   85.0  66.0  29.0  26.6
```

```
[113]: # Drop the target variable 'type' from X
X = data_dummy.drop(['type'], axis=1)

# Standardize the features
scale = StandardScaler().fit(X)
features = scale.transform(X)
features_scaled = pd.DataFrame(features, columns=["npreg", "ped", "age", "glu", "bp", "skin", "bmi"])

# Create the target variable 'y'
y = pd.DataFrame(data_dummy['type'])

# Add a constant term to the features
X_scaled = sm.add_constant(features_scaled)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=10)

# Ensure the indices are aligned
X_train.reset_index(drop=True, inplace=True)
y_train.reset_index(drop=True, inplace=True)
```

```
[114]: Log_Reg_Model_Std_Scalar=sm.Logit(y_train,X_train).fit()
print(Log_Reg_Model_Std_Scalar.summary())
y_pred_prob=Log_Reg_Model_Std_Scalar.predict(X_test)
y_pred=["0" if x<0.5 else "1" for x in y_pred_prob]
y_pred=np.array(y_pred,dtype=np.float32)
```

Optimization terminated successfully.

Current function value: 0.429189

Iterations 6

Logit Regression Results

=====

Dep. Variable:	type	No. Observations:	432
Model:	Logit	Df Residuals:	424
Method:	MLE	Df Model:	7
Date:	Sun, 28 Jan 2024	Pseudo R-squ.:	0.3257
Time:	16:34:04	Log-Likelihood:	-185.41
converged:	True	LL-Null:	-274.97
Covariance Type:	nonrobust	LLR p-value:	2.975e-35

	coef	std err	z	P> z	[0.025	0.975]
const	-0.9423	0.137	-6.901	0.000	-1.210	-0.675
npreg	0.7030	0.175	4.024	0.000	0.361	1.045
ped	0.4401	0.138	3.182	0.001	0.169	0.711
age	0.1710	0.175	0.976	0.329	-0.172	0.514
glu	1.1143	0.153	7.286	0.000	0.815	1.414
bp	-0.1328	0.140	-0.948	0.343	-0.408	0.142
skin	-0.0001	0.176	-0.001	1.000	-0.345	0.345
bmi	0.6937	0.179	3.876	0.000	0.343	1.044

```
[115]: plot_confusion_matrix(Log_Reg_Model_Std_Scalar)
test_report = get_test_report(Log_Reg_Model_Std_Scalar)
print(test_report)
plot_roc(Log_Reg_Model_Std_Scalar)
update_score_card(model_name = 'Log_Reg_Model_Std_Scalar')
```

	Actual:0	60	13
	Actual:1	9	27
		Predicted:0	Predicted:1


	precision	recall	f1-score	support
0	0.87	0.82	0.85	73
1	0.68	0.75	0.71	36

accuracy			0.80	109
macro avg	0.77	0.79	0.78	109
weighted avg	0.81	0.80	0.80	109

```
[115]:
```

	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286
5	Naive_Bayes_Model	0.847291	0.720000
6	svm_linear	0.838095	0.767442
7	svm_poly	0.849261	0.857143
8	Log_Reg_Model_Std_Scalar	0.851598	0.675000

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754
5	0.779141	0.779141	0.502880	0.666667
6	0.785276	0.785276	0.502832	0.653465
7	0.797546	0.797546	0.515362	0.645161
8	0.798165	0.798165	0.556255	0.710526



```
[116]: # Drop the target variable 'type' from X
X = data_dummy.drop(['type'], axis=1)
```

```

# Standardize the features
scale = StandardScaler().fit(X)
features = scale.transform(X)
features_scaled = pd.DataFrame(features, columns=["npreg", "ped", "age", "glu",
↳ "bp", "skin", "bmi"])

# Create the target variable 'y'
y = pd.DataFrame(data_dummy['type'])

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features_scaled, y,
↳ test_size=0.2, random_state=10)

# Ensure the indices are aligned
X_train.reset_index(drop=True, inplace=True)
y_train.reset_index(drop=True, inplace=True)

from sklearn.neighbors import KNeighborsClassifier

KN_Classifier = KNeighborsClassifier(n_neighbors=28, p=2, metric="euclidean")
↳ #SVC(kernel='poly', probability=True) # Specify 'probability=True' to
↳ enable probability estimates
KN_Classifier_st = KN_Classifier.fit(X_train, y_train)
y_pred_proba = KN_Classifier_st.predict_proba(X_test)[: ,1]
y_pred = KN_Classifier_st .predict(X_test)
plot_confusion_matrix(KN_Classifier_st)
test_report = get_test_report(KN_Classifier_st)
print(test_report)
plot_roc(KN_Classifier_st)
update_score_card(model_name = 'KN_Classifier_StandardScaler')

```

Actual:	Actual:0	65	8
	Actual:1	12	24
		Predicted:0	Predicted:1


	precision	recall	f1-score	support
0	0.84	0.89	0.87	73
1	0.75	0.67	0.71	36
accuracy			0.82	109
macro avg	0.80	0.78	0.79	109
weighted avg	0.81	0.82	0.81	109

[116]:

	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286
5	Naive_Bayes_Model	0.847291	0.720000
6	svm_linear	0.838095	0.767442
7	svm_poly	0.849261	0.857143
8	Log_Reg_Model_Std_Scalar	0.851598	0.675000
9	KN_Classifier_Standard_Scaler	0.847793	0.750000

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754
5	0.779141	0.779141	0.502880	0.666667
6	0.785276	0.785276	0.502832	0.653465
7	0.797546	0.797546	0.515362	0.645161
8	0.798165	0.798165	0.556255	0.710526
9	0.816514	0.816514	0.573218	0.705882





```
[117]: from sklearn.ensemble import RandomForestClassifier
#intantiate the regressor
rf_cls = RandomForestClassifier(n_estimators=100, random_state=10)

# fit the regressor with training dataset
rf_cls.fit(X_train, y_train)
```

```
[117]: RandomForestClassifier(random_state=10)
```

```
[118]: # predict the values on test dataset using predict()
y_pred = rf_cls.predict(X_test)
y_pred
```

```
[118]: array([0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0,
        1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1,
        1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0,
        0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0],
        dtype=int64)
```

```
[119]: plot_confusion_matrix(rf_cls)
test_report = get_test_report(rf_cls)
print(test_report)
plot_roc(rf_cls)
update_score_card(model_name = 'rf_cls')
```


Actual:	Actual:0	59	14
	Actual:1	13	23
		Predicted:0	Predicted:1

	precision	recall	f1-score	support
0	0.82	0.81	0.81	73
1	0.62	0.64	0.63	36
accuracy			0.75	109
macro avg	0.72	0.72	0.72	109
weighted avg	0.75	0.75	0.75	109

[119]:	Model	AUC Score	Precision Score \
0	Logistic_Regression with Full Model	0.864835	0.833333
1	Logistic Regression (SGD)	0.642674	0.636364
2	Log_Reg_Without_Intercept	0.744689	0.607143
3	Log_Reg_Backward_Model_Selection	0.863370	0.833333
4	decision_tree_grid	0.842611	0.714286
5	Naive_Bayes_Model	0.847291	0.720000
6	svm_linear	0.838095	0.767442
7	svm_poly	0.849261	0.857143
8	Log_Reg_Model_Std_Scalar	0.851598	0.675000
9	KN_Classifier_Standard_Scaler	0.847793	0.750000
10	rf_cls	0.847793	0.621622

	Recall Score	Accuracy Score	Kappa Score	f1-Score
0	0.788991	0.788991	0.498098	0.634921
1	0.697248	0.697248	0.270829	0.459016
2	0.697248	0.697248	0.297324	0.507463
3	0.788991	0.788991	0.498098	0.634921
4	0.791411	0.791411	0.541453	0.701754
5	0.779141	0.779141	0.502880	0.666667

6	0.785276	0.785276	0.502832	0.653465
7	0.797546	0.797546	0.515362	0.645161
8	0.798165	0.798165	0.556255	0.710526
9	0.816514	0.816514	0.573218	0.705882
10	0.752294	0.752294	0.443983	0.630137



```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

## 18 Multiple Regression Model Building for ped using PCA

```
[137]: from numpy.linalg import eig
        from sklearn.decomposition import PCA
```

```
[138]: data_1 = data_dummy.copy(deep = True)
        data_1.head()
```

```
[138]:
```

	npreg	ped	age	type	glu	bp	skin	bmi
0	6	0.627	50	1	148.0	72.0	35.0	33.6
1	1	0.351	31	0	85.0	66.0	29.0	26.6
3	1	0.167	21	0	89.0	66.0	23.0	28.1
4	0	2.288	33	1	137.0	40.0	35.0	43.1
6	3	0.248	26	1	78.0	50.0	32.0	31.0

```
[139]: for feature in ["type"]:  
        data_1[feature] = data_1[feature].astype('object')
```

```
[140]: data.dtypes
```

```
[140]: npreg      int64  
ped      float64  
age      int64  
type     int64  
glu      float64  
bp       float64  
skin     float64  
bmi      float64  
dtype: object
```

## 19 Compute Principal Components (from scratch)

## 20 Perform PCA with the following steps:

1. Filter the numerical variables
2. Scale the data to get variables on the same scale
3. Compute covariance matrix
4. Calculate eigenvalues and eigenvectors of the covariance matrix
5. Decide the number of principal components
6. Obtain principal components

```
[141]: df_num_features = data_1.select_dtypes(include=[np.number])  
data_num = df_num_features.drop('ped', axis=1)  
data_num.head()
```

```
[141]:   npreg  age  glu  bp  skin  bmi  
0      6  50 148.0 72.0 35.0 33.6  
1      1  31  85.0 66.0 29.0 26.6  
3      1  21  89.0 66.0 23.0 28.1  
4      0  33 137.0 40.0 35.0 43.1  
6      3  26  78.0 50.0 32.0 31.0
```

```
[142]: data_num_std = StandardScaler().fit_transform(data_num)  
print(data_num_std)
```

```
[[ 0.74901305  1.71809765  0.87972826  0.04375637  0.55855696  0.10282275]  
 [-0.7562472  -0.05200616 -1.16844217 -0.4460132  -0.01465704 -0.91866136]  
 [-0.7562472  -0.98363975 -1.03839961 -0.4460132  -0.58787104 -0.69977191]  
 ...  
 [-0.45519515 -0.4246596   0.03445158 -0.11950015 -0.20572838  0.56978691]  
 [ 0.447961   -0.14516952  0.00194093  0.04375637 -0.58787104 -0.97703189]  
 [-0.7562472  -0.79731303 -0.90835704 -0.11950015  0.17641429 -0.36414142]]
```

```
[143]: cov_mat = np.cov(data_num_std.T)
print(cov_mat[0:5])
```

```
[[1.00185185  0.64587984  0.12755508  0.20544545  0.1004247   0.01953158]
 [0.64587984  1.00185185  0.28023555  0.34778955  0.16712469  0.08173092]
 [0.12755508  0.28023555  1.00185185  0.21798405  0.22778961  0.24757349]
 [0.20544545  0.34778955  0.21798405  1.00185185  0.22714303  0.31008079]
 [0.1004247   0.16712469  0.22778961  0.22714303  1.00185185  0.649028   ]]
```

```
[144]: eig_values, eig_vector = np.linalg.eig(cov_mat)
print('Eigen values:', '\n', '\n', eig_values, "\n")
print('Eigen vectors:', '\n', '\n', eig_vector, '\n')
```

Eigen values:

```
[2.2991948  1.47162588 0.31010656 0.34945851 0.82499884 0.75572652]
```


Eigen vectors:

```
[[-0.36607157 -0.54358524 -0.44030593  0.47421335  0.2818529  -0.26892012]
 [-0.45281473 -0.4839029   0.61269163 -0.41773334  0.0524138  -0.0903184 ]
 [-0.35885404  0.0456002  -0.13289832  0.07190307 -0.8978727  -0.2003422 ]
 [-0.41488941 -0.01663631 -0.20873069 -0.02897317 -0.00391307  0.88496709]
 [-0.43110956  0.44434121 -0.40069046 -0.53421526  0.27934406 -0.30452184]
 [-0.4173554   0.52015173  0.45711743  0.55608014  0.18327011 -0.05905298]]
```

```
[145]: eig_values = list(eig_values)
eig_values.sort(reverse = True)
print(eig_values)
```

```
[2.299194798088914, 1.4716258844527035, 0.8249988427720018, 0.7557265225211692,
0.34945850631472203, 0.3101065569616022]
```

```
[146]: plt.plot(eig_values, 'bp')
plt.plot(eig_values)
plt.xlabel('Principal Components')
plt.ylabel('Percentage of explained variance')
plt.annotate(text = 'Elbow Point', xy=(3,0.8), xytext=(4, 2),
             ↪arrowprops=dict(facecolor='black', arrowstyle = 'simple'))
plt.title('Scree Plot')
plt.show()
```



Using the scree plot alone, it is challenging to make a decision on the number of principal components to select for extracting maximum variations in the data. Therefore, we employed a manual approach, calculating the percentage of variance explained by the first  $p$  eigenvalues. In our case, the first four eigenvalues exceed 90%. Consequently, we have selected the first four components for further analysis.

```
[147]: eigenvector = eig_vector[:,0:4]
       eigenvector
```

```
[147]: array([[ -0.36607157, -0.54358524, -0.44030593,  0.47421335],
              [ -0.45281473, -0.4839029 ,  0.61269163, -0.41773334],
              [ -0.35885404,  0.0456002 , -0.13289832,  0.07190307],
              [ -0.41488941, -0.01663631, -0.20873069, -0.02897317],
              [ -0.43110956,  0.44434121, -0.40069046, -0.53421526],
              [ -0.4173554 ,  0.52015173,  0.45711743,  0.55608014]])
```

```
[148]: pca_data = pd.DataFrame(data_num_std.dot(eigenvector), columns=[
      ↪ ['PC1', 'PC2', 'PC3', 'PC4'])
```

```
[149]: # Reset indices before concatenation
pca_data_reset = pca_data.reset_index(drop=True)
data_reset = data[['ped', 'type']].reset_index(drop=True)

# Concatenate along the index axis
data_pca = pd.concat([pca_data_reset, data_reset], axis=1)

# Now check for missing values
print("Missing values in 'data_pca':")
print(data_pca.isnull().sum())
```

```
Missing values in 'data_pca':
PC1      0
PC2      0
PC3      0
PC4      0
ped      0
type     0
dtype: int64
```

```
[150]: data2=data_pca.copy(True)
data2.head()
```

```
[150]:
```

	PC1	PC2	PC3	PC4	ped	type
0	-1.669733	-0.897484	0.420015	-0.541739	0.627	1
1	1.294463	-0.093966	0.135434	-0.911009	0.351	0
2	1.825418	0.221937	-0.122913	-0.084545	0.167	0
3	0.342151	1.599030	1.471433	0.084141	2.288	1
4	1.516979	0.276827	0.066443	-0.205483	0.248	1

```
[151]: data2.corr()
```


```
[151]:
```

	PC1	PC2	PC3	PC4	ped	\
PC1	1.000000e+00	-3.640388e-16	-3.574367e-16	1.123898e-17	-0.135099	
PC2	-3.640388e-16	1.000000e+00	-1.386431e-15	1.190572e-15	0.084609	
PC3	-3.574367e-16	-1.386431e-15	1.000000e+00	3.031616e-15	0.070958	
PC4	1.123898e-17	1.190572e-15	3.031616e-15	1.000000e+00	0.012846	
ped	-1.350988e-01	8.460946e-02	7.095778e-02	1.284589e-02	1.000000	
type	-4.891232e-01	-4.075420e-03	2.480845e-02	9.079427e-02	0.225474	


	type
PC1	-0.489123
PC2	-0.004075
PC3	0.024808
PC4	0.090794
ped	0.225474
type	1.000000

```
[152]: sns.heatmap(data2.corr(), cmap='Blues', annot=True)
plt.show()
```



```
[153]: data2.ped.hist(color = 'maroon')
plt.title('Distribution of Target Variable (ped)', fontsize = 15)
plt.xlabel('ped', fontsize = 15)
plt.ylabel('Frequency', fontsize = 15)

# display the plot
plt.show()
```



The distribution of the target variable is right-skewed. However, the assumption for the target variable in linear regression is a normal distribution. To address the right-skewed nature of the variable, a log transformation is applied to achieve normality.



```
[154]: import numpy as np
data2['log_ped'] = np.log(data2['ped'])

# display the top 5 rows of the data
data2.head()
```


```
[154]:      PC1      PC2      PC3      PC4  ped  type  log_ped
0 -1.669733 -0.897484  0.420015 -0.541739  0.627    1 -0.466809
1  1.294463 -0.093966  0.135434 -0.911009  0.351    0 -1.046969
2  1.825418  0.221937 -0.122913 -0.084545  0.167    0 -1.789761
3  0.342151  1.599030  1.471433  0.084141  2.288    1  0.827678
4  1.516979  0.276827  0.066443 -0.205483  0.248    1 -1.394327
```

```
[1]: #data_numeric = data2.select_dtypes(include=np.number)
#print(data_numeric.columns)
#data_categoric = data2.select_dtypes(include = object)
#dummy_variables = pd.get_dummies(data_categoric, drop_first = True)
```

```
[2]: #data_2 = pd.concat([data_numeric, dummy_variables], axis=1)
#data_2.head()
```

```
[317]: data_2.log_ped.hist(color = 'maroon')
plt.title('Distribution of Target Variable (log_ped)', fontsize = 15)
plt.xlabel('log_ped', fontsize = 15)
plt.ylabel('Frequency', fontsize = 15)

# display the plot
plt.show()
```



```
[318]: from sklearn.model_selection import train_test_split
X = data_2.drop(['ped', 'log_ped'], axis = 1)
import statsmodels
import statsmodels.api as sm
X=sm.add_constant(X)
y = pd.DataFrame(data_2[['ped', 'log_ped']])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ random_state = 1)
print("The shape of y_test is:", y_test.shape)
```

The shape of y\_test is: (109, 2)

```
[319]: from statsmodels.compat import lzip
import statsmodels.stats.api as sms
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
from statsmodels.tools.eval_measures import rmse
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
[320]: #Build model using sm.OLS().fit()
linreg_logmodel_full_pca = sm.OLS(y_train['log_ped'], X_train).fit()

# print the summary output
print(linreg_logmodel_full_pca.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	log_ped	R-squared:	0.071			
Model:	OLS	Adj. R-squared:	0.060			
Method:	Least Squares	F-statistic:	6.512			
Date:	Mon, 18 Dec 2023	Prob (F-statistic):	7.52e-06			
Time:	14:30:51	Log-Likelihood:	-398.47			
No. Observations:	432	AIC:	808.9			
Df Residuals:	426	BIC:	833.3			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	-1.0074	0.037	-27.174	0.000	-1.080	-0.935
PC1	-0.0009	0.022	-0.042	0.967	-0.045	0.043
PC2	0.0159	0.024	0.664	0.507	-0.031	0.063
PC3	0.1405	0.053	2.672	0.008	0.037	0.244
PC4	-0.0214	0.053	-0.405	0.685	-0.125	0.082

type_1	0.3080	0.072	4.307	0.000	0.167	0.449
=====						
Omnibus:		3.248	Durbin-Watson:			1.977
Prob(Omnibus):		0.197	Jarque-Bera (JB):			2.585
Skew:		0.058	Prob(JB):			0.275
Kurtosis:		2.639	Cond. No.			3.97
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The R-squared and adjusted R-squared values are considerably low, indicating poor performance of the model.

```
[321]: linreg_logmodel_full_predictions_pca = linreg_logmodel_full_pca.predict(X_test)
```

```
[322]: predicted_ped_pca = np.exp(linreg_logmodel_full_predictions_pca)

# extract the 'Property_Sale_Price' values from the test data
actual_ped= y_test['ped']
```

```
[323]: linreg_logmodel_full_rmse_pca = rmse(actual_ped, predicted_ped_pca)

# calculate R-squared using rsquared
linreg_logmodel_full_rsquared_pca = linreg_logmodel_full_pca.rsquared

# calculate Adjusted R-Squared using rsquared_adj
linreg_logmodel_full_rsquared_adj_pca = linreg_logmodel_full_pca.rsquared_adj
```

```
[324]: cols = ['Model', 'RMSE', 'R-Squared', 'Adj. R-Squared']
result_tabulation=pd.DataFrame(columns=cols)
linreg_logmodel_full_metrics = pd.Series({'Model': "LR model with log of target_
↪variable_PCA ",
                                           'RMSE':linreg_logmodel_full_rmse_pca,
                                           'R-Squared': linreg_logmodel_full_rsquared_pca,
                                           'Adj. R-Squared': linreg_logmodel_full_rsquared_adj_pca
                                           })
result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
↪ignore_index = True)
result_tabulation
```

```
[324]:
```

	Model	RMSE	R-Squared	\
0	LR model with log of target variable_PCA	0.40599	0.071001	
	Adj. R-Squared			
0	0.060097			

```
[327]: from sklearn.linear_model import SGDRegressor
SGD_model_pca=SGDRegressor(loss="squared_error",alpha=0.1,max_iter=1000)
SGD_model_pca.fit(X_train,y_train["log_ped"])
y_pred_SGD_pca=SGD_model_pca.predict(X_test)
X_train.shape
```

```
[327]: (432, 6)
```

```
[328]: r_square_SGD_pca=SGD_model_pca.score(X_train,y_train["log_ped"])
n=432
p=6
adj_R2_SGD_pca=1-(1-r_square_SGD_pca)*(n-1)/(n-p-1)
from sklearn.metrics import mean_squared_error
from math import sqrt
rmse_SGD_pca=sqrt(mean_squared_error(y_test["log_ped"],y_pred_SGD_pca))
```

```
[329]: linreg_logmodel_full_metrics = pd.Series({'Model': "SGD_Model_pca",'RMSE':
    ↪rmse_SGD_pca,'R-Squared':r_square_SGD_pca,'Adj. R-Squared':adj_R2_SGD_pca })
result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
    ↪ignore_index = True)
result_tabulation
```

```
[329]:
```

	Model	RMSE	R-Squared \
0	LR model with log of target variable_PCA	0.405990	0.071001
1	SGD_Model_pca	0.661159	0.053627

	Adj. R-Squared
0	0.060097
1	0.040267

```
[330]: data.head()
```

```
[330]:
```

	npreg	ped	age	type	glu	bp	skin	bmi
0	6	0.627	50	1	148.0	72.0	35.0	33.6
1	1	0.351	31	0	85.0	66.0	29.0	26.6
3	1	0.167	21	0	89.0	66.0	23.0	28.1
4	0	2.288	33	1	137.0	40.0	35.0	43.1
6	3	0.248	26	1	78.0	50.0	32.0	31.0

```
[331]: data.shape
```

```
[331]: (541, 8)
```

```
[332]: for feature in ["type"]:
    data[feature] = data[feature].astype('object')
```

```
[333]: import numpy as np
data['log_ped'] = np.log(data['ped'])

# display the top 5 rows of the data
data.head()
```

```
[333]:      npreg    ped  age type    glu    bp  skin    bmi  log_ped
0         6  0.627   50    1  148.0  72.0  35.0   33.6 -0.466809
1         1  0.351   31    0   85.0  66.0  29.0   26.6 -1.046969
3         1  0.167   21    0   89.0  66.0  23.0   28.1 -1.789761
4         0  2.288   33    1  137.0  40.0  35.0   43.1  0.827678
6         3  0.248   26    1   78.0  50.0  32.0   31.0 -1.394327
```

```
[334]: data_numeric = data.select_dtypes(include=np.number)
print(data_numeric.columns)
data_categorical = data.select_dtypes(include = object)
print(data_categorical.columns)
```

```
Index(['npreg', 'ped', 'age', 'glu', 'bp', 'skin', 'bmi', 'log_ped'],
      dtype='object')
Index(['type'], dtype='object')
```

```
[335]: dummy_variables = pd.get_dummies(data_categorical, drop_first = True)
data_dummy = pd.concat([data_numeric, dummy_variables], axis=1)
data_dummy
```

```
[335]:      npreg    ped  age    glu    bp  skin    bmi  log_ped  type_1
0         6  0.627   50  148.000000  72.0  35.0   33.6 -0.466809      1
1         1  0.351   31   85.000000  66.0  29.0   26.6 -1.046969      0
3         1  0.167   21   89.000000  66.0  23.0   28.1 -1.789761      0
4         0  2.288   33  137.000000  40.0  35.0   43.1  0.827678      1
6         3  0.248   26   78.000000  50.0  32.0   31.0 -1.394327      1
8         2  0.158   53  197.000000  70.0  45.0   30.5 -1.845160      1
13        1  0.398   59  189.000000  60.0  23.0   30.1 -0.921303      1
14        5  0.587   51  166.000000  72.0  19.0   25.8 -0.532730      1
16        0  0.551   31  118.000000  84.0  47.0   45.8 -0.596020      1
18        1  0.183   33  103.000000  30.0  38.0   43.3 -1.698269      0
19        1  0.529   32  115.000000  70.0  30.0   34.6 -0.636767      1
20        3  0.704   27  126.000000  88.0  41.0   39.3 -0.350977      0
23        9  0.263   29  119.000000  80.0  35.0   29.0 -1.335601      1
24       11  0.254   51  143.000000  94.0  33.0   36.6 -1.370421      1
25       10  0.205   41  125.000000  70.0  26.0   31.1 -1.584745      1
27        1  0.487   22   97.000000  66.0  15.0   23.2 -0.719491      0
28       13  0.245   57  145.000000  82.0  19.0   22.2 -1.406497      0
30        5  0.546   60  109.000000  75.0  26.0   36.0 -0.605136      0
31        3  0.851   28  158.000000  76.0  36.0   31.6 -0.161343      1
32        3  0.267   22   88.000000  58.0  11.0   24.8 -1.320507      0
```

34	10	0.512	45	122.000000	78.0	31.0	27.6	-0.669431	0
35	4	0.966	33	103.000000	60.0	33.0	24.0	-0.034591	0
37	9	0.665	46	102.000000	76.0	37.0	32.9	-0.407968	1
38	2	0.503	27	90.000000	68.0	42.0	38.2	-0.687165	1
39	4	1.390	56	111.000000	72.0	47.0	37.1	0.329304	1
40	3	0.271	26	180.000000	64.0	25.0	34.0	-1.305636	0
42	7	0.235	48	106.000000	92.0	18.0	22.7	-1.448170	0
43	9	0.721	54	171.000000	110.0	24.0	45.4	-0.327116	1
45	0	1.893	25	180.000000	66.0	39.0	42.0	0.638163	1
47	2	0.586	22	71.000000	70.0	27.0	28.0	-0.534435	0
48	7	0.344	31	103.000000	66.0	32.0	39.1	-1.067114	1
50	1	0.491	22	103.000000	80.0	11.0	19.4	-0.711311	0
51	1	0.526	26	101.000000	50.0	15.0	24.2	-0.642454	0
52	5	0.342	30	88.000000	66.0	21.0	24.4	-1.072945	0
53	8	0.467	58	176.000000	90.0	34.0	33.7	-0.761426	1
54	7	0.718	42	150.000000	66.0	42.0	34.7	-0.331286	0
55	1	0.248	21	73.000000	50.0	10.0	23.0	-1.394327	0
56	7	0.254	41	187.000000	68.0	39.0	37.7	-1.370421	1
57	0	0.962	31	100.000000	88.0	60.0	46.8	-0.038741	0
59	0	0.173	22	105.000000	64.0	41.0	41.5	-1.754464	0
63	2	0.699	24	141.000000	58.0	34.0	25.4	-0.358105	0
65	5	0.203	32	99.000000	74.0	27.0	29.0	-1.594549	0
66	0	0.855	38	109.000000	88.0	30.0	32.5	-0.156654	1
68	1	0.334	25	95.000000	66.0	13.0	19.6	-1.096614	0
69	4	0.189	27	146.000000	85.0	27.0	28.9	-1.666008	0
70	2	0.867	28	100.000000	66.0	20.0	32.9	-0.142716	1
71	5	0.411	26	139.000000	64.0	35.0	28.6	-0.889162	0
73	4	0.231	23	129.000000	86.0	20.0	35.1	-1.465338	0
74	1	0.396	22	79.000000	75.0	30.0	32.0	-0.926341	0
75	1	0.140	22	120.940299	48.0	20.0	24.7	-1.966113	0
77	5	0.370	27	95.000000	72.0	33.0	37.7	-0.994252	0
79	2	0.307	24	112.000000	66.0	22.0	25.0	-1.180908	0
80	3	0.140	22	113.000000	44.0	13.0	22.4	-1.966113	0
82	7	0.767	36	83.000000	78.0	26.0	29.3	-0.265268	0
83	0	0.237	22	101.000000	65.0	28.0	24.6	-1.439695	0
85	2	0.698	27	110.000000	74.0	29.0	32.4	-0.359536	0
86	13	0.178	45	106.000000	72.0	54.0	36.6	-1.725972	0
87	2	0.324	26	100.000000	68.0	25.0	38.5	-1.127012	0
88	15	0.153	43	136.000000	70.0	32.0	37.1	-1.877317	1
89	1	0.165	24	107.000000	68.0	19.0	26.5	-1.801810	0
91	4	0.443	34	123.000000	80.0	15.0	32.0	-0.814186	0
92	7	0.261	42	81.000000	78.0	40.0	46.7	-1.343235	0
94	2	0.761	21	142.000000	82.0	18.0	24.7	-0.273122	0
95	6	0.255	40	144.000000	72.0	27.0	33.9	-1.366492	0
96	2	0.130	24	92.000000	62.0	28.0	31.6	-2.040221	0
97	1	0.323	22	71.000000	48.0	18.0	20.4	-1.130103	0
98	6	0.356	23	93.000000	50.0	30.0	28.7	-1.032825	0

99	1	0.325	31	122.000000	90.0	51.0	49.7	-1.123930	1
103	1	0.283	24	81.000000	72.0	18.0	26.6	-1.262308	0
105	1	0.801	21	126.000000	56.0	29.0	28.7	-0.221894	0
107	4	0.287	37	144.000000	58.0	28.0	29.5	-1.248273	0
108	3	0.336	25	83.000000	58.0	31.0	34.3	-1.090644	0
109	0	0.247	24	95.000000	85.0	25.0	37.4	-1.398367	1
110	3	0.199	24	171.000000	72.0	33.0	33.3	-1.614450	1
111	8	0.543	46	155.000000	62.0	26.0	34.0	-0.610646	1
112	1	0.192	23	89.000000	76.0	34.0	31.2	-1.650260	0
114	7	0.588	39	160.000000	54.0	32.0	30.5	-0.531028	1
118	4	0.443	22	97.000000	60.0	23.0	28.2	-0.814186	0
119	4	0.223	21	99.000000	76.0	15.0	23.2	-1.500584	0
120	0	0.759	25	162.000000	76.0	56.0	53.2	-0.275754	1
121	6	0.260	24	111.000000	64.0	39.0	34.2	-1.347074	0
122	2	0.404	23	107.000000	74.0	30.0	33.6	-0.906340	0
125	1	0.496	26	88.000000	30.0	42.0	55.0	-0.701179	1
126	3	0.452	30	120.000000	70.0	30.0	42.9	-0.794073	0
127	1	0.261	23	118.000000	58.0	36.0	33.3	-1.343235	0
128	1	0.403	40	117.000000	88.0	24.0	34.5	-0.908819	1
130	4	0.361	33	173.000000	70.0	14.0	29.7	-1.018877	1
132	3	0.356	30	170.000000	64.0	37.0	34.5	-1.032825	1
133	8	0.457	39	84.000000	74.0	31.0	38.3	-0.783072	0
134	2	0.647	26	96.000000	68.0	13.0	21.1	-0.435409	0
135	2	0.088	31	125.000000	60.0	20.0	33.8	-2.430418	0
136	0	0.597	21	100.000000	70.0	26.0	30.8	-0.515838	0
137	0	0.532	22	93.000000	60.0	25.0	28.7	-0.631112	0
139	5	0.159	28	105.000000	72.0	29.0	36.9	-1.838851	0
141	5	0.286	38	106.000000	82.0	30.0	39.5	-1.251763	0
142	2	0.318	22	108.000000	52.0	26.0	32.5	-1.145704	0
144	4	0.237	23	154.000000	62.0	31.0	32.8	-1.439695	0
145	0	0.572	21	102.000000	75.0	23.0	32.8	-0.558616	0
146	9	0.096	41	57.000000	80.0	37.0	32.8	-2.343407	0
147	2	1.400	34	106.000000	64.0	35.0	30.5	0.336472	0
149	2	0.085	22	90.000000	70.0	17.0	27.3	-2.465104	0
150	1	0.399	24	136.000000	74.0	50.0	37.4	-0.918794	0
152	9	1.189	42	156.000000	86.0	28.0	34.3	0.173113	1
153	1	0.687	23	153.000000	82.0	42.0	40.6	-0.375421	0
155	7	0.337	36	152.000000	88.0	44.0	50.0	-1.087672	1
156	2	0.637	21	99.000000	52.0	15.0	24.6	-0.450986	0
157	1	0.833	23	109.000000	56.0	21.0	25.2	-0.182722	0
158	2	0.229	22	88.000000	74.0	19.0	29.0	-1.474033	0
159	17	0.817	47	163.000000	72.0	41.0	40.9	-0.202116	1
160	4	0.294	36	151.000000	90.0	38.0	29.7	-1.224176	0
161	7	0.204	45	102.000000	74.0	40.0	37.2	-1.589635	0
162	0	0.167	27	114.000000	80.0	34.0	44.2	-1.789761	0
163	2	0.368	21	100.000000	64.0	23.0	29.7	-0.999672	0
165	6	0.722	41	104.000000	74.0	18.0	29.9	-0.325730	1

166	3	0.256	22	148.000000	66.0	25.0	32.5	-1.362578	0
169	3	0.495	29	111.000000	90.0	12.0	28.4	-0.703198	0
171	6	0.542	29	134.000000	70.0	23.0	35.4	-0.612489	1
172	2	0.773	25	87.000000	72.0	23.0	28.9	-0.257476	0
173	1	0.678	23	79.000000	60.0	42.0	43.5	-0.388608	0
174	2	0.370	33	75.000000	64.0	24.0	29.7	-0.994252	0
175	8	0.719	36	179.000000	72.0	42.0	32.7	-0.329894	1
177	0	0.319	26	129.000000	110.0	46.0	67.1	-1.142564	1
181	0	0.725	23	119.000000	64.0	18.0	34.9	-0.321584	0
182	1	0.299	21	120.940299	74.0	20.0	27.7	-1.207312	0
185	7	0.745	41	194.000000	68.0	28.0	35.9	-0.294371	1
186	8	0.615	60	181.000000	68.0	36.0	30.1	-0.486133	1
187	1	1.321	33	128.000000	98.0	41.0	32.0	0.278389	1
188	8	0.640	31	109.000000	76.0	39.0	27.9	-0.446287	1
189	5	0.361	25	139.000000	80.0	35.0	31.6	-1.018877	1
191	9	0.374	40	123.000000	70.0	44.0	33.1	-0.983499	0
194	8	0.136	42	85.000000	55.0	20.0	24.4	-1.995100	0
195	5	0.395	29	158.000000	84.0	41.0	39.4	-0.928870	1
197	3	0.678	23	107.000000	62.0	13.0	22.9	-0.388608	1
198	4	0.905	26	109.000000	64.0	44.0	34.8	-0.099820	1
199	4	0.150	29	148.000000	60.0	27.0	30.9	-1.897120	1
200	0	0.874	21	113.000000	80.0	16.0	31.0	-0.134675	0
202	0	0.787	32	108.000000	68.0	20.0	27.3	-0.239527	0
203	2	0.235	27	99.000000	70.0	16.0	20.4	-1.448170	0
204	6	0.324	55	103.000000	72.0	32.0	37.7	-1.127012	0
205	5	0.407	27	111.000000	72.0	28.0	23.9	-0.898942	0
206	8	0.605	57	196.000000	76.0	29.0	37.5	-0.502527	1
208	1	0.289	21	96.000000	64.0	27.0	33.2	-1.241329	0
209	7	0.355	41	184.000000	84.0	33.0	35.5	-1.035637	1
210	2	0.290	25	81.000000	60.0	22.0	27.7	-1.237874	0
211	0	0.375	24	147.000000	85.0	54.0	42.8	-0.980829	0
212	7	0.164	60	179.000000	95.0	31.0	34.2	-1.807889	0
213	0	0.431	24	140.000000	65.0	26.0	42.6	-0.841647	1
214	9	0.260	36	112.000000	82.0	32.0	34.2	-1.347074	1
215	12	0.742	38	151.000000	70.0	40.0	41.8	-0.298406	1
216	5	0.514	25	109.000000	62.0	41.0	35.8	-0.665532	1
217	6	0.464	32	125.000000	68.0	30.0	30.0	-0.767871	0
218	5	1.224	32	85.000000	74.0	22.0	29.0	0.202124	1
220	0	1.072	21	177.000000	60.0	29.0	34.6	0.069526	1
223	7	0.687	61	142.000000	60.0	33.0	28.8	-0.375421	0
224	1	0.666	26	100.000000	66.0	15.0	23.6	-0.406466	0
225	1	0.101	22	87.000000	78.0	27.0	34.6	-2.292635	0
227	3	0.652	24	162.000000	52.0	38.0	37.2	-0.427711	1
228	4	2.329	31	197.000000	70.0	39.0	36.7	0.845439	0
229	0	0.089	24	117.000000	80.0	31.0	45.2	-2.419119	0
231	6	0.238	46	134.000000	80.0	37.0	46.2	-1.435485	1
232	1	0.583	22	79.000000	80.0	25.0	25.4	-0.539568	0



234	3	0.293	23	74.000000	68.0	28.0	29.7	-1.227583	0
236	7	0.586	51	181.000000	84.0	21.0	35.9	-0.534435	1
237	0	0.686	23	179.000000	90.0	27.0	44.1	-0.376878	1
238	9	0.831	32	164.000000	84.0	21.0	30.8	-0.185125	1
240	1	0.192	21	91.000000	64.0	24.0	29.2	-1.650260	0
241	4	0.446	22	91.000000	70.0	32.0	33.1	-0.807436	0
243	6	1.318	33	119.000000	50.0	22.0	27.1	0.276115	1
244	2	0.329	29	146.000000	76.0	35.0	38.2	-1.111698	0
245	9	1.213	49	184.000000	85.0	15.0	30.0	0.193097	1
247	0	0.427	23	165.000000	90.0	33.0	52.3	-0.850971	0
248	9	0.282	34	124.000000	70.0	33.0	35.4	-1.265848	0
249	1	0.143	23	111.000000	86.0	19.0	30.1	-1.944911	0
252	2	0.249	24	90.000000	80.0	14.0	24.4	-1.390302	0
253	0	0.238	25	86.000000	68.0	32.0	35.8	-1.435485	0
254	12	0.926	44	92.000000	62.0	7.0	27.6	-0.076881	1
255	1	0.543	21	113.000000	64.0	35.0	33.6	-0.610646	1
256	3	0.557	30	111.000000	56.0	39.0	30.1	-0.585190	0
257	2	0.092	25	114.000000	68.0	22.0	28.7	-2.385967	0
258	1	0.655	24	193.000000	50.0	16.0	25.9	-0.423120	0
259	11	1.353	51	155.000000	76.0	28.0	33.3	0.302324	1
260	3	0.299	34	191.000000	68.0	15.0	30.9	-1.207312	0
262	4	0.612	24	95.000000	70.0	32.0	32.1	-0.491023	0
263	3	0.200	63	142.000000	80.0	15.0	32.4	-1.609438	0
265	5	0.997	43	96.000000	74.0	18.0	33.6	-0.003005	0
267	2	1.101	24	128.000000	64.0	42.0	40.0	0.096219	0
270	10	1.136	38	101.000000	86.0	37.0	45.6	0.127513	1
271	2	0.128	21	108.000000	62.0	32.0	25.2	-2.055725	0
273	1	0.422	21	71.000000	78.0	50.0	33.2	-0.862750	0
275	2	0.677	25	100.000000	70.0	52.0	40.5	-0.390084	0
276	7	0.296	29	106.000000	60.0	24.0	26.5	-1.217396	1
277	0	0.454	23	104.000000	64.0	23.0	27.8	-0.789658	0
279	2	0.881	22	108.000000	62.0	10.0	25.3	-0.126698	0
281	10	0.280	39	129.000000	76.0	28.0	35.9	-1.272966	0
282	7	0.262	37	133.000000	88.0	15.0	32.4	-1.339411	0
285	7	0.647	51	136.000000	74.0	26.0	26.0	-0.435409	0
286	5	0.619	34	155.000000	84.0	44.0	38.7	-0.479650	0
287	1	0.808	29	119.000000	86.0	39.0	45.6	-0.213193	1
288	4	0.340	26	96.000000	56.0	17.0	20.8	-1.078810	0
289	5	0.263	33	108.000000	72.0	43.0	36.1	-1.335601	0
290	0	0.434	21	78.000000	88.0	29.0	36.9	-0.834711	0
291	0	0.757	25	107.000000	62.0	30.0	36.6	-0.278392	1
292	2	1.224	31	128.000000	78.0	37.0	43.3	0.202124	1
293	1	0.613	24	128.000000	48.0	45.0	40.5	-0.489390	1
295	6	0.692	28	151.000000	62.0	31.0	35.5	-0.368169	0
296	2	0.337	29	146.000000	70.0	38.0	28.0	-1.087672	1
297	0	0.520	24	126.000000	84.0	29.0	30.7	-0.653926	0
298	14	0.412	46	100.000000	78.0	25.0	36.6	-0.886732	1

301	2	0.422	25	144.000000	58.0	33.0	31.6	-0.862750	1
302	5	0.156	35	77.000000	82.0	41.0	35.8	-1.857899	0
305	2	0.215	29	120.000000	76.0	37.0	39.7	-1.537117	0
306	10	0.326	47	161.000000	68.0	23.0	25.5	-1.120858	1
307	0	0.143	21	137.000000	68.0	14.0	24.8	-1.944911	0
308	0	1.391	25	128.000000	68.0	19.0	30.5	0.330023	1
309	2	0.875	30	124.000000	68.0	28.0	32.9	-0.133531	1
310	6	0.313	41	80.000000	66.0	30.0	26.2	-1.161552	0
311	0	0.605	22	106.000000	70.0	37.0	39.4	-0.502527	0
312	2	0.433	27	155.000000	74.0	17.0	26.6	-0.837018	1
313	3	0.626	25	113.000000	50.0	10.0	29.5	-0.468405	0
314	7	1.127	43	109.000000	80.0	31.0	35.9	0.119559	1
315	2	0.315	26	112.000000	68.0	22.0	34.1	-1.155183	0
316	3	0.284	30	99.000000	80.0	11.0	19.3	-1.258781	0
318	3	0.150	28	115.000000	66.0	39.0	38.1	-1.897120	0
320	4	0.527	31	129.000000	60.0	12.0	27.5	-0.640555	0
321	3	0.197	25	112.000000	74.0	30.0	31.6	-1.624552	1
322	0	0.254	36	124.000000	70.0	20.0	27.4	-1.370421	1
323	13	0.731	43	152.000000	90.0	33.0	26.8	-0.313342	1
324	2	0.148	21	112.000000	75.0	32.0	35.7	-1.910543	0
325	1	0.123	24	157.000000	72.0	21.0	25.6	-2.095571	0
326	1	0.692	30	122.000000	64.0	32.0	35.1	-0.368169	1
328	2	0.127	23	102.000000	86.0	36.0	45.5	-2.063568	1
329	6	0.122	37	105.000000	70.0	32.0	30.8	-2.103734	0
330	8	1.476	46	118.000000	72.0	19.0	23.1	0.389336	0
331	2	0.166	25	87.000000	58.0	16.0	32.7	-1.795767	0
334	1	0.260	22	95.000000	60.0	18.0	23.9	-1.347074	0
335	0	0.259	26	165.000000	76.0	43.0	47.9	-1.350927	0
338	9	0.893	33	152.000000	78.0	34.0	34.2	-0.113169	1
340	1	0.472	22	130.000000	70.0	13.0	25.9	-0.750776	0
341	1	0.673	36	95.000000	74.0	21.0	25.9	-0.396010	0
342	1	0.389	22	120.940299	68.0	35.0	32.0	-0.944176	0
345	8	0.349	49	126.000000	88.0	36.0	38.5	-1.052683	0
346	1	0.654	22	139.000000	46.0	19.0	28.7	-0.424648	0
348	3	0.279	26	99.000000	62.0	19.0	21.8	-1.276543	0
349	5	0.346	37	120.940299	80.0	32.0	41.0	-1.061317	1
352	3	0.243	46	61.000000	82.0	28.0	34.4	-1.414694	0
353	1	0.580	24	90.000000	62.0	12.0	27.2	-0.544727	0
356	1	0.962	28	125.000000	50.0	40.0	33.3	-0.038741	1
357	13	0.569	44	129.000000	72.0	30.0	39.9	-0.563875	1
358	12	0.378	48	88.000000	74.0	40.0	35.3	-0.972861	0
359	1	0.875	29	196.000000	76.0	36.0	36.5	-0.133531	1
360	5	0.583	29	189.000000	64.0	33.0	31.2	-0.539568	1
362	5	0.305	65	103.000000	108.0	37.0	39.2	-1.187444	0
364	4	0.385	30	147.000000	74.0	25.0	34.9	-0.954512	0
365	5	0.499	30	99.000000	54.0	28.0	34.0	-0.695149	0
367	0	0.252	21	101.000000	64.0	17.0	21.0	-1.378326	0

368	3	0.306	22	81.000000	86.0	16.0	27.5	-1.184170	0
369	1	0.234	45	133.000000	102.0	28.0	32.8	-1.452434	1
370	3	2.137	25	173.000000	82.0	48.0	38.4	0.759403	1
371	0	1.731	21	118.000000	64.0	23.0	32.8	0.548699	0
372	0	0.545	21	84.000000	64.0	22.0	35.8	-0.606969	0
373	2	0.225	25	105.000000	58.0	40.0	34.9	-1.491655	0
374	2	0.816	28	122.000000	52.0	43.0	36.2	-0.203341	0
375	12	0.528	58	140.000000	82.0	43.0	39.2	-0.638659	1
376	0	0.299	22	98.000000	82.0	15.0	25.2	-1.207312	0
377	1	0.509	22	87.000000	60.0	37.0	37.2	-0.675307	0
379	0	1.021	35	93.000000	100.0	39.0	43.4	0.020783	0
380	1	0.821	24	107.000000	72.0	30.0	30.8	-0.197232	0
381	0	0.236	22	105.000000	68.0	22.0	20.0	-1.443923	0
382	1	0.947	21	109.000000	60.0	8.0	25.4	-0.054456	0
383	1	1.268	25	90.000000	62.0	18.0	25.1	0.237441	0
384	1	0.221	25	125.000000	70.0	24.0	24.3	-1.509593	0
385	1	0.205	24	119.000000	54.0	13.0	22.3	-1.584745	0
386	5	0.660	35	116.000000	74.0	29.0	32.3	-0.415515	1
387	8	0.239	45	105.000000	100.0	36.0	43.3	-1.431292	1
388	5	0.452	58	144.000000	82.0	26.0	32.0	-0.794073	1
389	3	0.949	28	100.000000	68.0	23.0	31.6	-0.052346	0
390	1	0.444	42	100.000000	66.0	29.0	32.0	-0.811931	0
392	1	0.389	21	131.000000	64.0	14.0	23.7	-0.944176	0
393	4	0.463	37	116.000000	72.0	12.0	22.1	-0.770028	0
395	2	1.600	25	127.000000	58.0	24.0	27.7	0.470004	0
396	3	0.944	39	96.000000	56.0	34.0	24.7	-0.057629	0
397	0	0.196	22	131.000000	66.0	40.0	34.3	-1.629641	1
399	3	0.241	25	193.000000	70.0	31.0	34.9	-1.422958	1
402	5	0.286	35	136.000000	84.0	41.0	35.0	-1.251763	1
403	9	0.280	38	72.000000	78.0	25.0	31.6	-1.272966	0
405	2	0.520	26	123.000000	48.0	32.0	42.1	-0.653926	0
409	1	0.702	28	172.000000	68.0	49.0	42.4	-0.353822	1
410	6	0.674	28	102.000000	90.0	39.0	35.7	-0.394525	0
411	1	0.528	25	112.000000	72.0	30.0	34.4	-0.638659	0
412	1	1.076	22	143.000000	84.0	23.0	42.4	0.073250	0
413	1	0.256	21	143.000000	74.0	22.0	26.2	-1.362578	0
414	0	0.534	21	138.000000	60.0	35.0	34.6	-0.627359	1
415	3	0.258	22	173.000000	84.0	33.0	35.7	-1.354796	1
416	1	1.095	22	97.000000	68.0	21.0	27.2	0.090754	0
417	4	0.554	37	144.000000	82.0	32.0	38.5	-0.590591	1
419	3	0.219	28	129.000000	64.0	29.0	26.4	-1.518684	1
420	1	0.507	26	119.000000	88.0	41.0	45.3	-0.679244	0
421	2	0.561	21	94.000000	68.0	18.0	26.0	-0.578034	0
422	0	0.496	21	102.000000	64.0	46.0	40.6	-0.701179	0
423	2	0.421	21	115.000000	64.0	22.0	30.8	-0.865122	0
424	8	0.516	36	151.000000	78.0	32.0	42.9	-0.661649	1
425	4	0.264	31	184.000000	78.0	39.0	37.0	-1.331806	1

427	1	0.328	38	181.000000	64.0	30.0	34.1	-1.114742	1
428	0	0.284	26	135.000000	94.0	46.0	40.6	-1.258781	0
429	1	0.233	43	95.000000	82.0	25.0	35.0	-1.456717	1
431	3	0.551	38	89.000000	74.0	16.0	30.4	-0.596020	0
432	1	0.527	22	80.000000	74.0	11.0	30.0	-0.640555	0
434	1	1.138	36	90.000000	68.0	8.0	24.5	0.129272	0
436	12	0.244	41	140.000000	85.0	33.0	37.4	-1.410587	0
438	1	0.147	21	97.000000	70.0	15.0	18.2	-1.917323	0
440	0	0.435	41	189.000000	104.0	25.0	34.3	-0.832409	1
441	2	0.497	22	83.000000	66.0	23.0	32.2	-0.699165	0
442	4	0.230	24	117.000000	64.0	27.0	33.2	-1.469676	0
444	4	0.380	30	117.000000	62.0	12.0	29.7	-0.967584	1
445	0	2.420	25	180.000000	78.0	63.0	59.4	0.883768	1
446	1	0.658	28	100.000000	72.0	12.0	25.3	-0.418550	0
447	0	0.330	26	95.000000	80.0	45.0	36.5	-1.108663	0
448	0	0.510	22	104.000000	64.0	37.0	33.6	-0.673345	1
449	0	0.285	26	120.000000	74.0	18.0	30.5	-1.255266	0
450	1	0.415	23	82.000000	64.0	13.0	21.2	-0.879477	0
452	0	0.381	25	91.000000	68.0	32.0	39.9	-0.964956	0
454	2	0.498	24	100.000000	54.0	28.0	37.8	-0.697155	0
455	14	0.212	38	175.000000	62.0	30.0	33.6	-1.551169	1
457	5	0.364	24	86.000000	68.0	28.0	30.2	-1.010601	0
458	10	1.001	51	148.000000	84.0	48.0	37.6	0.001000	1
459	9	0.460	81	134.000000	74.0	33.0	25.9	-0.776529	0
460	9	0.733	48	120.000000	72.0	22.0	20.8	-0.310610	0
462	8	0.705	39	74.000000	70.0	40.0	35.3	-0.349557	0
463	5	0.258	37	88.000000	78.0	30.0	27.6	-1.354796	0
465	0	0.452	21	124.000000	56.0	13.0	21.8	-0.794073	0
466	0	0.269	22	74.000000	52.0	10.0	27.8	-1.313044	0
467	0	0.600	25	97.000000	64.0	36.0	36.8	-0.510826	0
469	6	0.571	27	154.000000	78.0	41.0	46.1	-0.560366	0
470	1	0.607	28	144.000000	82.0	40.0	41.3	-0.499226	0
471	0	0.170	22	137.000000	70.0	38.0	33.2	-1.771957	0
472	0	0.259	22	119.000000	66.0	27.0	38.8	-1.350927	0
475	0	0.231	59	137.000000	84.0	27.0	27.3	-1.465338	0
476	2	0.711	29	105.000000	80.0	45.0	33.7	-0.341083	1
477	7	0.466	31	114.000000	76.0	17.0	23.8	-0.763570	0
478	8	0.162	39	126.000000	74.0	38.0	25.9	-1.820159	0
479	4	0.419	63	132.000000	86.0	31.0	28.0	-0.869884	0
480	3	0.344	35	158.000000	70.0	30.0	35.5	-1.067114	1
481	0	0.197	29	123.000000	88.0	37.0	35.2	-1.624552	0
482	4	0.306	28	85.000000	58.0	22.0	27.8	-1.184170	0
483	0	0.233	23	84.000000	82.0	31.0	38.2	-1.456717	0
485	0	0.365	24	135.000000	68.0	42.0	42.3	-1.007858	1
486	1	0.536	21	139.000000	62.0	41.0	40.7	-0.623621	0
487	0	1.159	58	173.000000	78.0	32.0	46.5	0.147558	0
488	4	0.294	28	99.000000	72.0	17.0	25.6	-1.224176	0

490	2	0.629	24	83.000000	65.0	28.0	36.8	-0.463624	0
491	2	0.292	42	89.000000	90.0	30.0	33.5	-1.231001	0
492	4	0.145	33	99.000000	68.0	38.0	32.8	-1.931022	0
493	4	1.144	45	125.000000	70.0	18.0	28.9	0.134531	1
497	2	0.547	25	81.000000	72.0	15.0	30.1	-0.603306	0
498	7	0.163	55	195.000000	70.0	33.0	25.1	-1.814005	1
499	6	0.839	39	154.000000	74.0	32.0	29.3	-0.175545	0
500	2	0.313	21	117.000000	90.0	19.0	25.2	-1.161552	0
501	3	0.267	28	84.000000	72.0	32.0	37.2	-1.320507	0
502	6	0.727	41	120.940299	68.0	41.0	39.0	-0.318829	1
503	7	0.738	41	94.000000	64.0	25.0	33.3	-0.303811	0
504	3	0.238	40	96.000000	78.0	39.0	37.3	-1.435485	0
506	0	0.314	35	180.000000	90.0	26.0	36.5	-1.158362	1
507	1	0.692	21	130.000000	60.0	23.0	28.6	-0.368169	0
508	2	0.968	21	84.000000	50.0	23.0	30.4	-0.032523	0
510	12	0.297	46	84.000000	72.0	31.0	29.7	-1.214023	1
511	0	0.207	21	139.000000	62.0	17.0	22.1	-1.575036	0
514	3	0.154	24	99.000000	54.0	19.0	25.6	-1.870803	0
515	3	0.268	28	163.000000	70.0	18.0	31.6	-1.316768	1
516	9	0.771	53	145.000000	88.0	34.0	30.3	-0.260067	1
519	6	0.582	60	129.000000	90.0	7.0	19.6	-0.541285	0
520	2	0.187	25	68.000000	70.0	32.0	25.0	-1.676647	0
521	3	0.305	26	124.000000	80.0	33.0	33.2	-1.187444	0
525	3	0.444	21	87.000000	60.0	18.0	21.8	-0.811931	0
526	1	0.299	21	97.000000	64.0	19.0	18.2	-1.207312	0
527	3	0.107	24	116.000000	74.0	15.0	26.3	-2.234926	0
528	0	0.493	22	117.000000	66.0	31.0	30.8	-0.707246	0
530	2	0.717	22	122.000000	60.0	18.0	29.8	-0.332679	0
532	1	0.917	29	86.000000	66.0	52.0	41.3	-0.086648	0
534	1	1.251	24	77.000000	56.0	30.0	33.3	0.223943	0
538	0	0.804	23	127.000000	80.0	37.0	36.3	-0.218156	0
539	3	0.968	32	129.000000	92.0	49.0	36.4	-0.032523	1
540	8	0.661	43	100.000000	74.0	40.0	39.4	-0.414001	1
541	3	0.549	27	128.000000	72.0	25.0	32.4	-0.599657	1
542	10	0.825	56	90.000000	85.0	32.0	34.9	-0.192372	1
543	4	0.159	25	84.000000	90.0	23.0	39.5	-1.838851	0
544	1	0.365	29	88.000000	78.0	29.0	32.0	-1.007858	0
545	8	0.423	37	186.000000	90.0	35.0	34.5	-0.860383	1
546	5	1.034	53	187.000000	76.0	27.0	43.6	0.033435	1
547	4	0.160	28	131.000000	68.0	21.0	33.1	-1.832581	0
548	1	0.341	50	164.000000	82.0	43.0	32.8	-1.075873	0
549	4	0.680	37	189.000000	110.0	31.0	28.5	-0.385662	0
550	1	0.204	21	116.000000	70.0	28.0	27.4	-1.589635	0
551	3	0.591	25	84.000000	68.0	30.0	31.9	-0.525939	0
553	1	0.422	23	88.000000	62.0	24.0	29.9	-0.862750	0
554	1	0.471	28	84.000000	64.0	23.0	36.9	-0.752897	0
555	7	0.161	37	124.000000	70.0	33.0	25.5	-1.826351	0

556	1	0.218	30	97.000000	70.0	40.0	38.1	-1.523260	0
558	11	0.126	42	103.000000	68.0	40.0	46.2	-2.071473	0
561	0	0.502	28	198.000000	66.0	32.0	41.3	-0.689155	1
562	1	0.401	24	87.000000	68.0	34.0	37.6	-0.913794	0
563	6	0.497	32	99.000000	60.0	19.0	26.9	-0.699165	0
565	2	0.748	22	95.000000	54.0	14.0	26.1	-0.290352	0
566	1	0.412	21	99.000000	72.0	30.0	38.6	-0.886732	0
567	6	0.085	46	92.000000	62.0	32.0	32.0	-2.465104	0
568	4	0.338	37	154.000000	72.0	29.0	31.3	-1.084709	0
569	0	0.203	33	121.000000	66.0	30.0	34.3	-1.594549	1
572	3	0.430	22	111.000000	58.0	31.0	29.5	-0.843970	0
573	2	0.198	22	98.000000	60.0	17.0	34.7	-1.619488	0
574	1	0.892	23	143.000000	86.0	30.0	30.1	-0.114289	0
575	1	0.280	25	119.000000	44.0	47.0	35.5	-1.272966	0
576	6	0.813	35	108.000000	44.0	20.0	24.0	-0.207024	0
579	2	0.575	62	197.000000	70.0	99.0	34.7	-0.553385	1
580	0	0.371	21	151.000000	90.0	46.0	42.1	-0.991553	1
581	6	0.206	27	109.000000	60.0	27.0	25.0	-1.579879	0
582	12	0.259	62	121.000000	78.0	17.0	26.5	-1.350927	0
584	8	0.687	52	124.000000	76.0	24.0	28.7	-0.375421	1
585	1	0.417	22	93.000000	56.0	11.0	22.5	-0.874669	0
588	3	1.154	52	176.000000	86.0	27.0	33.3	0.143234	1
590	11	0.925	45	111.000000	84.0	40.0	46.8	-0.077962	1
591	2	0.175	24	112.000000	78.0	50.0	39.4	-1.742969	0
593	2	1.699	25	82.000000	52.0	22.0	28.5	0.530040	0
594	6	0.733	34	123.000000	72.0	45.0	33.6	-0.310610	0
595	0	0.682	22	188.000000	82.0	14.0	32.0	-0.382726	1
597	1	0.559	21	89.000000	24.0	19.0	27.8	-0.581606	0
599	1	0.407	26	109.000000	38.0	18.0	23.1	-0.898942	0
600	1	0.400	24	108.000000	88.0	19.0	27.1	-0.916291	0
602	1	0.100	30	124.000000	74.0	36.0	27.8	-2.302585	0
603	7	0.692	54	150.000000	78.0	29.0	35.2	-0.368169	1
605	1	0.514	21	124.000000	60.0	32.0	35.8	-0.665532	0
606	1	1.258	22	181.000000	78.0	42.0	40.0	0.229523	1
607	1	0.482	25	92.000000	62.0	25.0	19.5	-0.729811	0
608	0	0.270	27	152.000000	82.0	39.0	41.5	-1.309333	0
609	1	0.138	23	111.000000	62.0	13.0	24.0	-1.980502	0
610	3	0.292	24	106.000000	54.0	21.0	30.9	-1.231001	0
611	3	0.593	36	174.000000	58.0	22.0	32.9	-0.522561	1
612	7	0.787	40	168.000000	88.0	42.0	38.2	-0.239527	1
613	6	0.878	26	105.000000	80.0	28.0	32.5	-0.130109	0
614	11	0.557	50	138.000000	74.0	26.0	36.1	-0.585190	1
617	2	0.257	23	68.000000	62.0	13.0	20.1	-1.358679	0
618	9	1.282	50	112.000000	82.0	24.0	28.2	0.248421	1
620	2	0.246	28	112.000000	86.0	42.0	38.4	-1.402424	0
621	2	1.698	28	92.000000	76.0	20.0	24.2	0.529451	0
623	0	0.347	21	94.000000	70.0	27.0	43.5	-1.058430	0

625	4	0.362	29	90.000000	88.0	47.0	37.7	-1.016111	0
629	4	0.148	21	94.000000	65.0	22.0	24.7	-1.910543	0
631	0	0.238	24	102.000000	78.0	40.0	34.5	-1.435485	0
633	1	0.115	22	128.000000	82.0	17.0	27.5	-2.162823	0
637	2	0.649	23	94.000000	76.0	18.0	31.6	-0.432323	0
638	7	0.871	32	97.000000	76.0	32.0	40.9	-0.138113	1
639	1	0.149	28	100.000000	74.0	12.0	19.5	-1.903809	0
640	0	0.695	27	102.000000	86.0	17.0	29.3	-0.363843	0
644	3	0.730	27	103.000000	72.0	30.0	27.6	-0.314711	0
645	2	0.134	30	157.000000	74.0	35.0	39.4	-2.009915	0
646	1	0.447	33	167.000000	74.0	17.0	23.4	-0.805197	1
647	0	0.455	22	179.000000	50.0	36.0	37.8	-0.787458	1
648	11	0.260	42	136.000000	84.0	35.0	28.3	-1.347074	1
649	0	0.133	23	107.000000	60.0	25.0	26.4	-2.017406	0
650	1	0.234	23	91.000000	54.0	25.0	25.2	-1.452434	0
651	1	0.466	27	117.000000	60.0	23.0	33.8	-0.763570	0
652	5	0.269	28	123.000000	74.0	40.0	34.1	-1.313044	0
654	1	0.142	22	106.000000	70.0	28.0	34.2	-1.951928	0
655	2	0.240	25	155.000000	52.0	27.0	38.7	-1.427116	1
656	2	0.155	22	101.000000	58.0	35.0	21.8	-1.864330	0
657	1	1.162	41	120.000000	80.0	48.0	38.9	0.150143	0
659	3	1.292	27	80.000000	82.0	31.0	34.2	0.256191	1
661	1	1.394	22	199.000000	76.0	43.0	42.9	0.332177	1
662	8	0.165	43	167.000000	106.0	46.0	37.6	-1.801810	1
663	9	0.637	40	145.000000	80.0	46.0	37.9	-0.450986	1
664	6	0.245	40	115.000000	60.0	39.0	33.7	-1.406497	1
665	1	0.217	24	112.000000	80.0	45.0	34.8	-1.527858	0
666	4	0.235	70	145.000000	82.0	18.0	32.5	-1.448170	1
667	10	0.141	40	111.000000	70.0	27.0	27.5	-1.958995	1
668	6	0.430	43	98.000000	58.0	33.0	34.0	-0.843970	0
669	9	0.164	45	154.000000	78.0	30.0	30.9	-1.807889	0
670	6	0.631	49	165.000000	68.0	26.0	33.6	-0.460449	0
671	1	0.551	21	99.000000	58.0	10.0	25.4	-0.596020	0
672	10	0.285	47	68.000000	106.0	23.0	35.5	-1.255266	0
673	3	0.880	22	123.000000	100.0	35.0	57.3	-0.127833	0
679	2	0.614	23	101.000000	58.0	17.0	24.2	-0.487760	0
680	2	0.332	22	56.000000	56.0	28.0	24.2	-1.102620	0
681	0	0.364	26	162.000000	76.0	36.0	49.6	-1.010601	1
682	0	0.366	22	95.000000	64.0	39.0	44.6	-1.005122	0
685	2	0.591	25	129.000000	74.0	26.0	33.2	-0.525939	0
687	1	0.181	29	107.000000	50.0	19.0	28.3	-1.709258	0
688	1	0.828	23	140.000000	74.0	26.0	24.1	-0.188742	0
689	1	0.335	46	144.000000	82.0	46.0	46.1	-1.093625	1
692	2	0.886	23	121.000000	70.0	32.0	39.1	-0.121038	0
693	7	0.439	43	129.000000	68.0	49.0	38.5	-0.823256	1
695	7	0.128	43	142.000000	90.0	24.0	30.4	-2.055725	1
696	3	0.268	31	169.000000	74.0	19.0	29.9	-1.316768	1

698	4	0.598	28	127.000000	88.0	11.0	34.5	-0.514165	0
700	2	0.483	26	122.000000	76.0	27.0	35.9	-0.727739	0
701	6	0.565	49	125.000000	78.0	31.0	27.6	-0.570930	1
702	1	0.905	52	168.000000	88.0	29.0	35.0	-0.099820	1
704	4	0.118	27	110.000000	76.0	20.0	28.4	-2.137071	0
705	6	0.177	28	80.000000	80.0	36.0	39.8	-1.731606	0
707	2	0.176	22	127.000000	46.0	21.0	34.4	-1.737271	0
709	2	0.674	23	93.000000	64.0	32.0	38.0	-0.394525	1
710	3	0.295	24	158.000000	64.0	13.0	31.2	-1.220780	0
711	5	0.439	40	126.000000	78.0	27.0	29.6	-0.823256	0
712	10	0.441	38	129.000000	62.0	36.0	41.2	-0.818710	1
713	0	0.352	21	134.000000	58.0	20.0	26.4	-1.044124	0
715	7	0.826	34	187.000000	50.0	33.0	33.9	-0.191161	1
716	3	0.970	31	173.000000	78.0	39.0	33.8	-0.030459	1
717	10	0.595	56	94.000000	72.0	18.0	23.1	-0.519194	0
718	1	0.415	24	108.000000	60.0	46.0	35.5	-0.879477	0
719	5	0.378	52	97.000000	76.0	27.0	35.6	-0.972861	1
720	4	0.317	34	83.000000	86.0	19.0	29.3	-1.148854	0
721	1	0.289	21	114.000000	66.0	36.0	38.1	-1.241329	0
722	1	0.349	42	149.000000	68.0	29.0	29.3	-1.052683	1
723	5	0.251	42	117.000000	86.0	30.0	39.1	-1.382302	0
725	4	0.236	38	112.000000	78.0	40.0	39.4	-1.443923	0
726	1	0.496	25	116.000000	78.0	29.0	36.1	-0.701179	0
727	0	0.433	22	141.000000	84.0	26.0	32.4	-0.837018	0
730	3	0.323	34	130.000000	78.0	23.0	28.4	-1.130103	1
732	2	0.646	24	174.000000	88.0	37.0	44.5	-0.436956	1
733	2	0.426	22	106.000000	56.0	27.0	29.0	-0.853316	0
735	4	0.284	28	95.000000	60.0	32.0	35.4	-1.258781	0
736	0	0.515	21	126.000000	86.0	27.0	27.4	-0.663588	0
737	8	0.600	42	65.000000	72.0	23.0	32.0	-0.510826	0
738	2	0.453	21	99.000000	60.0	17.0	36.6	-0.791863	0
740	11	0.785	48	120.000000	80.0	37.0	42.3	-0.242072	1
741	3	0.400	26	102.000000	44.0	20.0	30.8	-0.916291	0
742	1	0.219	22	109.000000	58.0	18.0	28.5	-1.518684	0
744	13	1.174	39	153.000000	88.0	37.0	40.6	0.160417	0
745	12	0.488	46	100.000000	84.0	33.0	30.0	-0.717440	0
746	1	0.358	27	147.000000	94.0	41.0	49.3	-1.027222	1
747	1	1.096	32	81.000000	74.0	41.0	46.3	0.091667	0
748	3	0.408	36	187.000000	70.0	22.0	36.4	-0.896488	1
751	1	0.261	28	121.000000	78.0	39.0	39.0	-1.343235	0
752	3	0.223	25	108.000000	62.0	24.0	26.0	-1.500584	0
753	0	0.222	26	181.000000	88.0	44.0	43.3	-1.505078	1
754	8	0.443	45	154.000000	78.0	32.0	32.4	-0.814186	1
755	1	1.057	37	128.000000	88.0	39.0	36.5	0.055435	1
756	7	0.391	39	137.000000	90.0	41.0	32.0	-0.939048	0
760	2	0.766	22	88.000000	58.0	26.0	28.4	-0.266573	0
761	9	0.403	43	170.000000	74.0	31.0	44.0	-0.908819	1



763	10	0.171	63	101.000000	76.0	48.0	32.9	-1.766092	0
764	2	0.340	27	122.000000	70.0	27.0	36.8	-1.078810	0
765	5	0.245	30	121.000000	72.0	23.0	26.2	-1.406497	0
767	1	0.315	23	93.000000	70.0	31.0	30.4	-1.155183	0

```
[336]: from sklearn.model_selection import train_test_split
X = data_dummy.drop(['ped', 'log_ped'], axis = 1)
import statsmodels
import statsmodels.api as sm
X=sm.add_constant(X)
y = pd.DataFrame(data_dummy[['ped', 'log_ped']])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
random_state = 1)
print("The shape of y_test is:", y_test.shape)
```

The shape of y\_test is: (109, 2)

```
[337]: #Build model using sm.OLS().fit()
linreg_logmodel_full_intercept = sm.OLS(y_train['log_ped'], X_train).fit()

# print the summary output
print(linreg_logmodel_full_intercept.summary())
```

OLS Regression Results

```
=====
```

Dep. Variable:	log_ped	R-squared:	0.073
Model:	OLS	Adj. R-squared:	0.057
Method:	Least Squares	F-statistic:	4.738
Date:	Mon, 18 Dec 2023	Prob (F-statistic):	3.85e-05
Time:	14:35:02	Log-Likelihood:	-398.11
No. Observations:	432	AIC:	812.2
Df Residuals:	424	BIC:	844.8
Df Model:	7		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
const	-1.0605	0.227	-4.669	0.000	-1.507	-0.614
npreg	-0.0220	0.011	-1.916	0.056	-0.045	0.001
age	0.0086	0.004	2.232	0.026	0.001	0.016
glu	-0.0005	0.001	-0.433	0.665	-0.003	0.002
bp	-0.0044	0.003	-1.619	0.106	-0.010	0.001
skin	-0.0024	0.004	-0.607	0.544	-0.010	0.005
bmi	0.0093	0.006	1.515	0.131	-0.003	0.021
type_1	0.2986	0.077	3.894	0.000	0.148	0.449
-----	-----	-----	-----	-----	-----	-----

```
=====
```

Omnibus:	3.716	Durbin-Watson:	1.981
----------	-------	----------------	-------

Prob(Omnibus):	0.156	Jarque-Bera (JB):	2.867
Skew:	0.062	Prob(JB):	0.238
Kurtosis:	2.620	Cond. No.	1.19e+03

=====

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.19e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[338]: linreg_logmodel_full_intercept_predictions = linreg_logmodel_full_intercept.
        ↪predict(X_test)
        predicted_ped = np.exp(linreg_logmodel_full_intercept_predictions)

        # extract the 'ped' values from the test data
        actual_ped= y_test['ped']
        linreg_logmodel_full_intercept_rmse = rmse(actual_ped, predicted_ped)

        # calculate R-squared using rsquared
        linreg_logmodel_full_intercept_rsquared = linreg_logmodel_full_intercept.
        ↪rsquared

        # calculate Adjusted R-Squared using rsquared_adj
        linreg_logmodel_full_intercept_rsquared_adj = linreg_logmodel_full_intercept.
        ↪rsquared_adj
```

```
[339]: linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg full model with log of target_intercept",
        ↪'RMSE':linreg_logmodel_full_intercept_rmse,
        ↪'R-Squared': linreg_logmodel_full_intercept_rsquared,
        ↪'Adj. R-Squared':linreg_logmodel_full_intercept_rsquared_adj
        })
        result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
        ↪ignore_index = True)
        result_tabulation
```

```
[339]:
```

	Model	RMSE	R-Squared \
0	LR model with log of target variable_PCA	0.405990	0.071001
1	SGD_Model_pca	0.661159	0.053627
2	Linreg full model with log of target_intercept	0.405092	0.072553

	Adj. R-Squared
0	0.060097
1	0.040267
2	0.057242

## 21 Linear Regression Model without Intercept

```
[340]: X = data_dummy.drop(['ped', 'log_ped'], axis = 1)
y = pd.DataFrame(data_dummy[['ped', 'log_ped']])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ random_state = 1)
```

```
[341]: #Build model using sm.OLS().fit()
linreg_logmodel_full= sm.OLS(y_train['log_ped'], X_train).fit()

# print the summary output
print(linreg_logmodel_full.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          log_ped    R-squared (uncentered):
0.681
Model:                  OLS      Adj. R-squared (uncentered):
0.676
Method:                 Least Squares    F-statistic:
129.7
Date:                   Mon, 18 Dec 2023    Prob (F-statistic):
2.67e-101
Time:                   14:35:25    Log-Likelihood:
-408.94
No. Observations:      432    AIC:
831.9
Df Residuals:          425    BIC:
860.4
Df Model:               7
Covariance Type:       nonrobust
=====
=====
```

	coef	std err	t	P> t	[0.025	0.975]
npreg	-0.0205	0.012	-1.749	0.081	-0.044	0.003
age	0.0060	0.004	1.520	0.129	-0.002	0.014
glu	-0.0027	0.001	-2.496	0.013	-0.005	-0.001
bp	-0.0100	0.002	-4.073	0.000	-0.015	-0.005
skin	-0.0007	0.004	-0.180	0.857	-0.009	0.007
bmi	-0.0021	0.006	-0.361	0.718	-0.013	0.009
type_1	0.4210	0.074	5.703	0.000	0.276	0.566

```
=====
=====
Omnibus:                1.869    Durbin-Watson:                1.978
Prob(Omnibus):          0.393    Jarque-Bera (JB):          1.793
Skew:                   0.085    Prob(JB):                  0.408
Kurtosis:               2.734    Cond. No.                  376.
=====
```

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The model exhibits an R-squared of 0.681 and an adjusted R-squared of 0.676, indicating a performance level of 68%. However, certain features, namely age, skin, and BMI, are found to be insignificant. To address this, feature selection techniques such as Lasso regression, k-best features, and backward elimination methods should be employed to identify and retain the most relevant features.

```
[342]: linreg_logmodel_full_predictions = linreg_logmodel_full.predict(X_test)
        predicted_ped = np.exp(linreg_logmodel_full_predictions)

        # extract the 'ped' values from the test data
        actual_ped= y_test['ped']
        linreg_logmodel_full_rmse = rmse(actual_ped, predicted_ped)

        # calculate R-squared using rsquared
        linreg_logmodel_full_rsquared = linreg_logmodel_full.rsquared

        # calculate Adjusted R-Squared using rsquared_adj
        linreg_logmodel_full_rsquared_adj = linreg_logmodel_full.rsquared_adj

        linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg full model with log of target",
                                                  'RMSE':linreg_logmodel_full_rmse,
                                                  'R-Squared': linreg_logmodel_full_rsquared,
                                                  'Adj. R-Squared': linreg_logmodel_full_rsquared_adj
                                                  })
        result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics, ignore_index = True)
        result_tabulation
```

```
[342]:
```

	Model	RMSE	R-Squared	\
0	LR model with log of target variable_PCA	0.405990	0.071001	
1	SGD_Model_pca	0.661159	0.053627	
2	Linreg full model with log of target_intercept	0.405092	0.072553	
3	Linreg full model with log of target	0.414520	0.681068	
	Adj. R-Squared			
0	0.060097			
1	0.040267			
2	0.057242			
3	0.675815			

```
[ ]:
```

```
[343]: data.head(1)
```

```
[343]:    npreg    ped  age type    glu    bp  skin    bmi  log_ped
0      6  0.627   50    1  148.0  72.0  35.0   33.6 -0.466809
```

## 22 Recursive Feature Elimination (RFE):

```
[344]: from sklearn.feature_selection import RFE
X = data_dummy.drop(['ped', 'log_ped'], axis = 1)
y = pd.DataFrame(data_dummy[['ped']])

model = LinearRegression()
rfe = RFE(model, n_features_to_select=4) # Specify the number of features you
    want to keep
fit = rfe.fit(X, y)
selected_features = X.columns[fit.support_].tolist()
selected_features
```

```
[344]: ['npreg', 'age', 'bmi', 'type_1']
```

The RFE technique has identified the features npreg, age, bmi, and type. Using these selected features, we need to construct a regression model and assess both its performance and the significance of the chosen features.

```
[345]: X = data_dummy.drop(['ped', 'log_ped', 'glu', 'bp', 'skin'], axis = 1)
y = pd.DataFrame(data_dummy[['ped', 'log_ped']])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
    random_state = 1)
linreg_logmodel_RFE = sm.OLS(y_train['log_ped'], X_train).fit()
print(linreg_logmodel_RFE.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          log_ped    R-squared (uncentered):
0.656
Model:                OLS    Adj. R-squared (uncentered):
0.653
Method:              Least Squares    F-statistic:
204.4
Date:                Mon, 18 Dec 2023    Prob (F-statistic):
7.31e-98
Time:                14:36:50    Log-Likelihood:
-425.02
No. Observations:    432    AIC:
```

858.0

Df Residuals: 428 BIC:

874.3

Df Model: 4

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
npreg	-0.0101	0.012	-0.846	0.398	-0.033	0.013
age	-0.0066	0.003	-1.966	0.050	-0.013	-7.93e-07
bmi	-0.0229	0.003	-8.530	0.000	-0.028	-0.018
type_1	0.4208	0.071	5.928	0.000	0.281	0.560
Omnibus:	1.483	Durbin-Watson:	1.932			
Prob(Omnibus):	0.476	Jarque-Bera (JB):	1.402			
Skew:	0.033	Prob(JB):	0.496			
Kurtosis:	2.729	Cond. No.	107.			

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The model demonstrates an R-squared of 0.656 and an adjusted R-squared of 0.653, representing a performance level of 66%. However, it was observed that the feature npreg is deemed insignificant in this context.

```
[346]: linreg_logmodel_RFE_predictions = linreg_logmodel_RFE.predict(X_test)
predicted_ped = np.exp(linreg_logmodel_RFE_predictions)

# extract the 'ped' values from the test data
actual_ped= y_test['ped']
linreg_logmodel_RFE_rmse = rmse(actual_ped, predicted_ped)

# calculate R-squared using rsquared
linreg_logmodel_RFE_rsquared = linreg_logmodel_RFE.rsquared

# calculate Adjusted R-Squared using rsquared_adj
linreg_logmodel_RFE_rsquared_adj = linreg_logmodel_RFE.rsquared_adj

linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg model with RFE",
      'RMSE':linreg_logmodel_RFE_rmse,
      'R-Squared': linreg_logmodel_RFE_rsquared,
      'Adj. R-Squared': linreg_logmodel_RFE_rsquared_adj
    })
```

```
result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
↪ ignore_index = True)
result_tabulation
```

```
[346]:
```

	Model	RMSE	R-Squared	\
0	LR model with log of target variable_PCA	0.405990	0.071001	
1	SGD_Model_pca	0.661159	0.053627	
2	Linreg full model with log of target_intercept	0.405092	0.072553	
3	Linreg full model with log of target	0.414520	0.681068	
4	Linreg model with RFE	0.411422	0.656409	

	Adj. R-Squared
0	0.060097
1	0.040267
2	0.057242
3	0.675815
4	0.653198

```
[347]: from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LinearRegression
from sklearn.datasets import make_regression

# Create a linear regression model
model = LinearRegression()

# Define the number of folds
num_folds = 4

# Create a k-fold cross-validation object
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)

# Perform k-fold cross-validation
cross_val_scores = cross_val_score(model, X, y, cv=kf,
↪ scoring='neg_mean_squared_error')
# 'cv' parameter takes the k-fold cross-validation object
# 'scoring' parameter defines the metric to evaluate (here, negative mean
↪ squared error)

# Display the cross-validation scores
print("Cross-Validation Scores:", cross_val_scores)

# Calculate the mean and standard deviation of cross-validation scores
print("Mean CV Score:", np.mean(cross_val_scores))
print("Std CV Score:", np.std(cross_val_scores))
```

Cross-Validation Scores: [-0.21820245 -0.23192636 -0.32524801 -0.24987971]

Mean CV Score: -0.2563141330277254  
Std CV Score: 0.04135379399276952

```
[348]: data.head(1)
```

```
[348]:   npreg   ped  age type   glu   bp  skin   bmi  log_ped
0      6  0.627   50    1 148.0  72.0  35.0  33.6 -0.466809
```

Similarly, we applied Lasso and k-best techniques to select the optimal features, and subsequently, regression models were employed. The tables below provide the model summaries for the Lasso and k-best features models.

## 23 LASSO (L1 Regularization):

```
[349]: X = data_dummy.drop(['ped', 'log_ped'], axis = 1)
import statsmodels
import statsmodels.api as sm
y = pd.DataFrame(data_dummy[['log_ped']])
```

```
[350]: from sklearn.linear_model import Lasso

lasso = Lasso(alpha=0.15) # You can experiment with different alpha values
model=lasso.fit(X, y)
selected_features = X.columns[model.coef_[0] != 0].tolist()
selected_features
```

```
[350]: []
```

```
[351]: X = data_dummy.drop(['ped', 'log_ped', 'npreg', 'type_1', 'age'], axis = 1)
y = pd.DataFrame(data_dummy[['ped', 'log_ped']])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ random_state = 1)
linreg_logmodel_Lasso= sm.OLS(y_train['log_ped'], X_train).fit()
print(linreg_logmodel_Lasso.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable:          log_ped   R-squared (uncentered):
0.654
Model:                  OLS      Adj. R-squared (uncentered):
0.651
Method:                 Least Squares   F-statistic:
202.3
Date:                   Mon, 18 Dec 2023   Prob (F-statistic):
3.15e-97
Time:                   14:37:29   Log-Likelihood:
```



```

-426.50
No. Observations:          432    AIC:
861.0
Df Residuals:              428    BIC:
877.3
Df Model:                  4
Covariance Type:          nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
glu          -0.0001      0.001      -0.109      0.914      -0.002      0.002
bp           -0.0118      0.002      -5.320      0.000      -0.016     -0.007
skin         0.0002      0.004       0.045      0.964      -0.008      0.008
bmi          -0.0008      0.006      -0.133      0.894      -0.012      0.011
=====
Omnibus:              4.541    Durbin-Watson:              1.951
Prob(Omnibus):         0.103    Jarque-Bera (JB):              3.267
Skew:                  0.049    Prob(JB):              0.195
Kurtosis:              2.586    Cond. No.              32.2
=====

```

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

[352]: linreg_logmodel_Lasso_predictions = linreg_logmodel_Lasso.predict(X_test)
       predicted_ped = np.exp(linreg_logmodel_RFE_predictions)

       # extract the 'ped' values from the test data
       actual_ped= y_test['ped']
       linreg_logmodel_Lasso_rmse = rmse(actual_ped, predicted_ped)

       # calculate R-squared using rsquared
       linreg_logmodel_Lasso_rsquared = linreg_logmodel_Lasso.rsquared

       # calculate Adjusted R-Squared using rsquared_adj
       linreg_logmodel_Lasso_rsquared_adj = linreg_logmodel_Lasso.rsquared_adj

       linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg model with Lasso",
                                                'RMSE':linreg_logmodel_Lasso_rmse,
                                                'R-Squared': linreg_logmodel_Lasso_rsquared,
                                                'Adj. R-Squared': linreg_logmodel_Lasso_rsquared_adj
                                                })
       result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
       ignore_index = True)

```

result\_tabulation

```
[352]:
```

	Model	RMSE	R-Squared	\
0	LR model with log of target variable_PCA	0.405990	0.071001	
1	SGD_Model_pca	0.661159	0.053627	
2	Linreg full model with log of target_intercept	0.405092	0.072553	
3	Linreg full model with log of target	0.414520	0.681068	
4	Linreg model with RFE	0.411422	0.656409	
5	Linreg model with Lasso	0.411422	0.654051	

	Adj. R-Squared
0	0.060097
1	0.040267
2	0.057242
3	0.675815
4	0.653198
5	0.650817

```
[353]: from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LinearRegression
from sklearn.datasets import make_regression

# Create a linear regression model
model = LinearRegression()

# Define the number of folds
num_folds = 4

# Create a k-fold cross-validation object
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)

# Perform k-fold cross-validation
cross_val_scores = cross_val_score(model, X, y, cv=kf,
    ↳scoring='neg_mean_squared_error')
# 'cv' parameter takes the k-fold cross-validation object
# 'scoring' parameter defines the metric to evaluate (here, negative mean
    ↳squared error)

# Display the cross-validation scores
print("Cross-Validation Scores:", cross_val_scores)

# Calculate the mean and standard deviation of cross-validation scores
print("Mean CV Score:", np.mean(cross_val_scores))
print("Std CV Score:", np.std(cross_val_scores))
```

Cross-Validation Scores: [-0.2377052 -0.24040804 -0.33661958 -0.23284152]

Mean CV Score: -0.2618935837578776  
Std CV Score: 0.043228181752897915

```
[354]: X.head(1)
```

```
[354]:      glu    bp  skin    bmi  
0  148.0  72.0  35.0  33.6
```

## 24 SelectKBest with f\_regression:

```
[355]: from sklearn.feature_selection import SelectKBest, f_regression  
X = data_dummy.drop(['ped', 'log_ped'], axis = 1)  
y = pd.DataFrame(data_dummy[['ped']])  
k_best = SelectKBest(score_func=f_regression, k=3) # Specify the number of  
↳ features you want to keep  
fit = k_best.fit(X, y)  
selected_features = X.columns[fit.get_support()].tolist()  
selected_features
```

```
[355]: ['glu', 'bmi', 'type_1']
```

```
[356]: X = data_dummy.drop(['ped', 'log_ped', 'npreg', 'bp', 'age', 'skin'], axis = 1)  
y = pd.DataFrame(data_dummy[['ped', 'log_ped']])  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,  
↳ random_state = 1)  
linreg_logmodel_Kbest = sm.OLS(y_train['log_ped'], X_train).fit()  
print(linreg_logmodel_Kbest.summary())
```

### OLS Regression Results

```
=====
```

Dep. Variable:	log_ped	R-squared (uncentered):
0.665		
Model:	OLS	Adj. R-squared (uncentered):
0.663		
Method:	Least Squares	F-statistic:
284.0		
Date:	Mon, 18 Dec 2023	Prob (F-statistic):
1.63e-101		
Time:	14:37:49	Log-Likelihood:
-419.46		
No. Observations:	432	AIC:
844.9		
Df Residuals:	429	BIC:
857.1		
Df Model:	3	
Covariance Type:	nonrobust	

	coef	std err	t	P> t	[0.025	0.975]
glu	-0.0045	0.001	-4.787	0.000	-0.006	-0.003
bmi	-0.0147	0.003	-4.451	0.000	-0.021	-0.008
type_1	0.4710	0.071	6.598	0.000	0.331	0.611
Omnibus:		0.448	Durbin-Watson:			1.956
Prob(Omnibus):		0.799	Jarque-Bera (JB):			0.518
Skew:		0.074	Prob(JB):			0.772
Kurtosis:		2.916	Cond. No.			299.

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

From the two tables above, it is evident that the K-best features selection model outperforms the Lasso features regression model. In the K-best model, all features are significant at any level of significance, whereas in the regression model, certain features do not exhibit significance.

```
[357]: linreg_logmodel_Kbest_predictions = linreg_logmodel_Kbest.predict(X_test)
predicted_ped = np.exp(linreg_logmodel_RFE_predictions)

# extract the 'ped' values from the test data
actual_ped= y_test['ped']
linreg_logmodel_Kbest_rmse = rmse(actual_ped, predicted_ped)

# calculate R-squared using rsquared
linreg_logmodel_Kbest_rsquared = linreg_logmodel_Kbest.rsquared

# calculate Adjusted R-Squared using rsquared_adj
linreg_logmodel_Kbest_rsquared_adj = linreg_logmodel_Kbest.rsquared_adj

linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg model with K-Best",
                                           'RMSE':linreg_logmodel_Kbest_rmse,
                                           'R-Squared': linreg_logmodel_Kbest_rsquared,
                                           'Adj. R-Squared': linreg_logmodel_Kbest_rsquared_adj
                                           })
result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
ignore_index = True)
result_tabulation
```

```
[357]:
```

	Model	RMSE	R-Squared	\
0	LR model with log of target variable_PCA	0.405990	0.071001	
1	SGD_Model_pca	0.661159	0.053627	

2	Linreg full model with log of target_intercept	0.405092	0.072553
3	Linreg full model with log of target	0.414520	0.681068
4	Linreg model with RFE	0.411422	0.656409
5	Linreg model with Lasso	0.411422	0.654051
6	Linreg model with K-Best	0.411422	0.665144

	Adj. R-Squared
0	0.060097
1	0.040267
2	0.057242
3	0.675815
4	0.653198
5	0.650817
6	0.662803

From the above metrics, we can observe that the linear regression model for K-best features performs better compared to other models. In this model, all features are significant. However, in the standard linear regression model, some features are found to be insignificant. Therefore, we choose the linear regression model with K-best features for analysis and deployment.

```
[ ]: # Drop the target variable 'type' from X
X = data.drop(['ped', 'type'], axis=1)

# Standardize the features
scale = StandardScaler().fit(X)
features = scale.transform(X)
features_scaled = pd.DataFrame(features, columns=["npreg", "age", "glu", "bp", "
↪skin", "bmi"])
db= pd.DataFrame(data['type'], columns=["type"])
scaled =features_scaled.reset_index(drop=True)
data_reset = data['type'].reset_index(drop=True)

# Concatenate along the index axis
data_5 = pd.concat([scaled, data_reset], axis=1)

# Now check for missing values
print("Missing values in 'data_pca':")
print(data_5.isnull().sum())

# Create the target variable 'y'
y =np.log(data['ped'])

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data_5, y, test_size=0.2,
↪random_state=10)
```

```
# Ensure the indices are aligned
X_train.reset_index(drop=True, inplace=True)
y_train.reset_index(drop=True, inplace=True)
data_5.head()
```

```
[ ]: linreg_logmodel_scaled= sm.OLS(y_train, X_train).fit()
print(linreg_logmodel_scaled.summary())
```

```
[ ]: linreg_logmodel_scaled_predictions =linreg_logmodel_scaled.predict(X_test)
predicted_ped = np.exp(linreg_logmodel_scaled_predictions)

# extract the 'ped' values from the test data
actual_ped= y_test['ped']
linreg_logmodel_scaled_rmse = rmse(actual_ped, predicted_ped)

# calculate R-squared using rsquared
linreg_logmodel_scaled_rsquared =linreg_logmodel_scaled.rsquared

# calculate Adjusted R-Squared using rsquared_adj
linreg_logmodel_scaled_rsquared_adj =linreg_logmodel_scaled.rsquared_adj

linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg model with Scaled_
↳Features",
                                           'RMSE':linreg_logmodel_scaled_rmse,
                                           'R-Squared':linreg_logmodel_scaled_rsquared,
                                           'Adj. R-Squared':linreg_logmodel_scaled_rsquared_adj
                                           })
result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
↳ignore_index = True)
result_tabulation
```

From the above metrics, we can observe that the linear regression model for K-best features performs better compared to other models. In this model, all features are significant. However, in the standard linear regression model, some features are found to be insignificant. Therefore, we choose the linear regression model with K-best features for analysis and deployment.

Performing k-fold cross-validation for linear regression involves splitting dataset into k folds, training the model on k-1 folds, and testing on the remaining fold. This process is repeated k times, each time using a different fold as the test set, and then the performance metrics are averaged over the k folds.

```
[358]: from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LinearRegression
from sklearn.datasets import make_regression
```

```

# Create a linear regression model
model = LinearRegression()

# Define the number of folds
num_folds = 4

# Create a k-fold cross-validation object
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)

# Perform k-fold cross-validation
cross_val_scores = cross_val_score(model, X, y, cv=kf,
    ↳scoring='neg_mean_squared_error')
# 'cv' parameter takes the k-fold cross-validation object
# 'scoring' parameter defines the metric to evaluate (here, negative mean
    ↳squared error)

# Display the cross-validation scores
print("Cross-Validation Scores:", cross_val_scores)

# Calculate the mean and standard deviation of cross-validation scores
print("Mean CV Score:", np.mean(cross_val_scores))
print("Std CV Score:", np.std(cross_val_scores))

```

```

Cross-Validation Scores: [-0.21852189 -0.22947818 -0.32575763 -0.25001804]
Mean CV Score: -0.2559439342203844
Std CV Score: 0.04186261537915351

```

In this case, the mean squared error values are relatively low, indicating that the model is performing well on the cross-validated data. The standard deviation is also relatively small, suggesting consistency across folds.

This is the final model selected under multiple linear regression. This model performs better than other models because this model contains all features (glucose, insulin, bmi, age and types) significant features at 5% level of significance. Also, R<sup>2</sup> and adjusted R<sup>2</sup> is almost similar to linear regression with lasso and recursive feature selection method but in these two models some of the features are insignificant. Also, my cross-validation method, the mean squared error values are relatively low, indicating that the model is performing well on the cross-validated data. The standard deviation is also relatively small, suggesting consistency across folds.

```

[188]: # Get predictions on the training set
       predictions = linreg_logmodel_Kbest.predict(X_train)

```

```

[189]: residuals = y_train['log_ped'] - predictions

```

```

[393]: data.head()

```

```

[393]:   npreg   ped  age type   glu   bp  skin   bmi  log_ped
0      6  0.627   50    1  148.0  72.0  35.0  33.6 -0.466809

```

1	1	0.351	31	0	85.0	66.0	29.0	26.6	-1.046969
2	1	0.167	21	0	89.0	66.0	23.0	28.1	-1.789761
3	0	2.288	33	1	137.0	40.0	35.0	43.1	0.827678
4	3	0.248	26	1	78.0	50.0	32.0	31.0	-1.394327

```
[400]: #data frame is data_dummy
X = data_dummy.drop(['ped','log_ped'], axis = 1)
y = pd.DataFrame(data_dummy[['log_ped']])

# Step 1: Start with the full model
model = sm.OLS(y, X).fit()

# Step 7: Stopping criterion (e.g., p-value threshold)
while model.pvalues.max() > 0.1:
    # Step 3: Evaluate predictor significance
    # Step 4: Remove the least significant predictor
    least_significant_predictor = model.pvalues.idxmax()
    X = X.drop(least_significant_predictor, axis=1)

    # Step 5: Fit the updated model
    model = sm.OLS(y,X).fit()

# Step 8: Final model
print(model.summary())
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:          log_ped    R-squared (uncentered):
0.661
Model:                  OLS      Adj. R-squared (uncentered):
0.659
Method:                 Least Squares    F-statistic:
349.9
Date:                   Mon, 18 Dec 2023    Prob (F-statistic):
5.60e-126
Time:                   12:32:43    Log-Likelihood:
-523.47
No. Observations:       541    AIC:
1053.
Df Residuals:           538    BIC:
1066.
Df Model:                3
Covariance Type:        nonrobust
=====
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----						



glu	-0.0017	0.001	-1.819	0.069	-0.003	0.000
bp	-0.0112	0.001	-7.649	0.000	-0.014	-0.008
type_1	0.4017	0.065	6.196	0.000	0.274	0.529

---

Omnibus:	3.206	Durbin-Watson:	2.004
Prob(Omnibus):	0.201	Jarque-Bera (JB):	2.581
Skew:	0.036	Prob(JB):	0.275
Kurtosis:	2.669	Cond. No.	339.

---

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[401]: X = data_dummy.drop(['ped', 'log_ped', 'npreg', 'age', 'skin', 'bmi'], axis = 1)
y = pd.DataFrame(data_dummy[['ped', 'log_ped']])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
random_state = 1)
linreg_logmodel_Backward= sm.OLS(y_train['log_ped'], X_train).fit()
print(linreg_logmodel_Backward.summary())
```

#### OLS Regression Results

---

Dep. Variable:	log_ped	R-squared (uncentered):	0.678
Model:	OLS	Adj. R-squared (uncentered):	0.676
Method:	Least Squares	F-statistic:	301.4
Date:	Mon, 18 Dec 2023	Prob (F-statistic):	3.22e-105
Time:	12:42:14	Log-Likelihood:	-410.86
No. Observations:	432	AIC:	827.7
Df Residuals:	429	BIC:	839.9
Df Model:	3		
Covariance Type:	nonrobust		

---

	coef	std err	t	P> t	[0.025	0.975]
glu	-0.0025	0.001	-2.443	0.015	-0.004	-0.000
bp	-0.0100	0.002	-6.168	0.000	-0.013	-0.007
type_1	0.4051	0.071	5.678	0.000	0.265	0.545

```
=====
Omnibus:                1.864    Durbin-Watson:                1.968
Prob(Omnibus):          0.394    Jarque-Bera (JB):        1.756
Skew:                   0.074    Prob(JB):                0.416
Kurtosis:               2.725    Cond. No.                340.
=====
```

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[402]: linreg_logmodel_Backward_predictions = linreg_logmodel_Backward.predict(X_test)
        predicted_ped = np.exp(linreg_logmodel_Backward_predictions)

        # extract the 'ped' values from the test data
        actual_ped= y_test['ped']
        linreg_logmodel_Backward_rmse = rmse(actual_ped, predicted_ped)

        # calculate R-squared using rsquared
        linreg_logmodel_Backward_rsquared =linreg_logmodel_Backward.rsquared

        # calculate Adjusted R-Squared using rsquared_adj
        linreg_logmodel_Backward_rsquared_adj =linreg_logmodel_Backward.rsquared_adj

        linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg model with Backward_
        ↪Selection",
                                                  'RMSE':linreg_logmodel_Backward_rmse,
                                                  'R-Squared':linreg_logmodel_Backward_rsquared,
                                                  'Adj. R-Squared':linreg_logmodel_Backward_rsquared_adj
                                                  })
        result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
        ↪ignore_index = True)
        result_tabulation
```

```
[402]:
```

	Model	RMSE	R-Squared \
0	LR model with log of target variable_PCA	0.382453	0.089738
1	SGD_Model_pca	0.626279	0.078270
2	Linreg full model with log of target_intercept	0.405092	0.072553
3	Linreg full model with log of target	0.414520	0.681068
4	Linreg model with RFE	0.411422	0.656409
5	Linreg model with Lasso	0.411422	0.654051
6	Linreg model with K-Best	0.411422	0.654051
7	Linreg model with Backward Selection	0.409188	0.678217

Adj. R-Squared

```

0      0.079055
1      0.056835
2      0.057242
3      0.675815
4      0.653198
5      0.650817
6      0.650817
7      0.675967

```

```
[40]: data_4.shape
```

```
[40]: (700, 7)
```

```
[44]:
```

```
Missing values in 'data_pca':
```

```

npreg    0
age      0
glu      0
bp       0
skin     0
bmi      0
type     0
dtype: int64

```

```
[44]:
```

	npreg	age	glu	bp	skin	bmi	type
0	0.749013	1.718098	0.879728	0.043756	0.558557	0.102823	1
1	-0.756247	-0.052006	-1.168442	-0.446013	-0.014657	-0.918661	0
2	-0.756247	-0.983640	-1.038400	-0.446013	-0.587871	-0.699772	0
3	-1.057299	0.134321	0.522111	-2.568348	0.558557	1.489123	1
4	-0.154143	-0.517823	-1.396017	-1.752065	0.271950	-0.276586	1

```
[45]: linreg_logmodel_scaled= sm.OLS(y_train, X_train).fit()
print(linreg_logmodel_scaled.summary())
```

#### OLS Regression Results

```

=====
=====
Dep. Variable:          ped    R-squared (uncentered):
0.239
Model:                OLS    Adj. R-squared (uncentered):
0.226
Method:               Least Squares    F-statistic:
19.05
Date:                 Thu, 21 Dec 2023    Prob (F-statistic):
3.81e-22
Time:                 12:35:58    Log-Likelihood:
-592.65

```

No. Observations: 432 AIC:  
1199.  
Df Residuals: 425 BIC:  
1228.  
Df Model: 7  
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
npreg	0.0925	0.064	1.452	0.147	-0.033	0.218
age	0.0679	0.068	0.995	0.320	-0.066	0.202
glu	0.2577	0.054	4.787	0.000	0.152	0.364
bp	-0.0472	0.051	-0.922	0.357	-0.148	0.053
skin	-0.0400	0.066	-0.609	0.543	-0.169	0.089
bmi	0.2114	0.064	3.307	0.001	0.086	0.337
type	-1.0162	0.090	-11.230	0.000	-1.194	-0.838
Omnibus:	1.167	Durbin-Watson:	1.356			
Prob(Omnibus):	0.558	Jarque-Bera (JB):	1.264			
Skew:	0.111	Prob(JB):	0.531			
Kurtosis:	2.855	Cond. No.	3.16			

Notes:

- [1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.  
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[ ]: linreg_logmodel_scaled_predictions = linreg_logmodel_scaled.predict(X_test)
predicted_ped = np.exp(linreg_logmodel_scaled_predictions)

# extract the 'ped' values from the test data
actual_ped= y_test['ped']
linreg_logmodel_scaled_rmse = rmse(actual_ped, predicted_ped)

# calculate R-squared using rsquared
linreg_logmodel_scaled_rsquared = linreg_logmodel_scaled.rsquared

# calculate Adjusted R-Squared using rsquared_adj
linreg_logmodel_scaled_rsquared_adj = linreg_logmodel_scaled.rsquared_adj

linreg_logmodel_full_metrics = pd.Series({'Model': "Linreg model with Scaled_
↳Features",
                                           'RMSE': linreg_logmodel_scaled_rmse,
                                           'R-Squared': linreg_logmodel_scaled_rsquared,
                                           'Adj. R-Squared': linreg_logmodel_scaled_rsquared_adj})
```

```

    })
result_tabulation = result_tabulation.append(linreg_logmodel_full_metrics,
↪ ignore_index = True)
result_tabulation

```

```

[359]: # Get predictions on the training set
predictions = linreg_logmodel_Kbest.predict(X_train)
residuals = y_train['log_ped'] - predictions

```

## 25 Residual Analysis

Create residual plots to visually assess the assumptions of linear regression. Common plots include the scatter plot of residuals against predicted values, a histogram of residuals, and a Q-Q plot:

## 26 Evaluate Homoscedasticity:


Check for homoscedasticity using a plot of residuals against predicted values. A cone-shaped pattern may indicate heteroscedasticity:

This type of plot is useful for identifying patterns or trends in the residuals. Ideally, we want to see a random scatter of points with no clear pattern, indicating that the residuals are homoscedastic. If we observe any systematic patterns (e.g., a funnel shape, curvature), it may suggest that the assumption of homoscedasticity is violated, and we might need to consider transformations or other model adjustments.

```

[191]: plt.figure(figsize=(7, 4))
plt.scatter(predictions, residuals)
plt.title('Residuals vs. Predicted Values')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.show()

```



In the scatter plot above, we can observe a random scatter of points with no clear pattern, suggesting that the residuals exhibit homoscedasticity. The absence of any systematic trend in the distribution of residuals indicates that the variance of the errors remains constant across different levels of predicted values

## 27 Check for Normality:

Assess the normality of residuals using a histogram and a Q-Q plot:

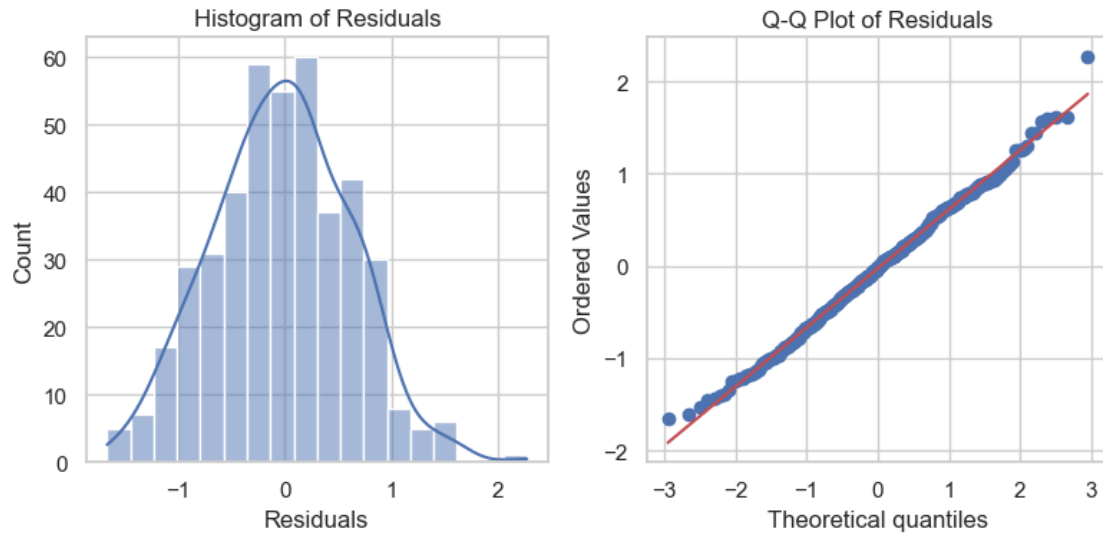
```
[371]: plt.figure(figsize=(8, 4))

# Histogram of residuals
plt.subplot(1, 2, 1)
sns.histplot(residuals, kde=True)
plt.title('Histogram of Residuals')
plt.xlabel('Residuals')

# Q-Q plot
plt.subplot(1, 2, 2)
stats.probplot(residuals, plot=plt)
```

```
plt.title('Q-Q Plot of Residuals')

plt.tight_layout()
plt.show()
```



```
[369]: import numpy as np
from scipy.stats import jarque_bera

# Perform Jarque-Bera test
statistic, p_value = jarque_bera(residuals)

# Display the results
print(f"Jarque-Bera statistic: {statistic}")
print(f"P-value: {p_value}")

# Check the null hypothesis
if p_value < 0.05:
    print("The Residuals are does not follows Normal distribution.")
else:
    print("The Residuals are follows Normal distribution.")
```

```
Jarque-Bera statistic: 0.5183927359564329
P-value: 0.7716714765970492
The Residuals are follows Normal distribution.
```

Based on the distribution plot, Q-Q plot, and the Jarque-Bera test, it has been confirmed that the residuals follow a normal distribution. This observation supports the assumption that the residuals are normally distributed, a key requirement in linear regression analysis.

## 28 Durbin-Watson Test:

The Durbin-Watson test checks for the presence of autocorrelation in the residuals. A value around 2 suggests no autocorrelation.

```
[195]: from statsmodels.stats.stattools import durbin_watson

# Calculate Durbin-Watson statistic
durbin_watson_stat = durbin_watson(linreg_logmodel_Kbest.resid)
print("Durbin-Watson Statistic:", durbin_watson_stat)
```

Durbin-Watson Statistic: 1.9485260033637735

In this case, a Durbin-Watson statistic of 1.95 is very close to 2, suggesting that there is no strong evidence of significant autocorrelation in the residuals. The value being close to 2 indicates that the residuals are approximately uncorrelated with each other.


## 29 Autocorrelation Plot:

Visualize the autocorrelation of residuals using an autocorrelation plot:

```
[373]: from statsmodels.graphics.tsaplots import plot_acf

# Autocorrelation plot
plt.figure(figsize=(5, 4))
plot_acf(linreg_logmodel_Kbest.resid)
plt.title('Autocorrelation of Residuals')
plt.show()
```

<Figure size 500x400 with 0 Axes>



The Durbin-Watson test statistic is approximately equal to 2, suggesting that the autocorrelation



function (ACF) of the error term is nearly zero. This indicates the absence of autocorrelation in the error term. Additionally, the Jarque-Bera (JB) test yields a p-value of 0.77. As the p-value is greater than the significance level (commonly 0.05), we accept the null hypothesis. This result implies that the residuals are normally distributed. Altogether, the Durbin-Watson test indicates no autocorrelation, and the JB test supports the normality assumption of the residuals.

## 30 Cluster Analysis

```
[84]: data.head()
```

```
[84]:
```

	npreg	ped	age	type	glu	bp	skin	bmi
0	6	0.627	50	1	148.0	72.0	35.0	33.6
1	1	0.351	31	0	85.0	66.0	29.0	26.6
3	1	0.167	21	0	89.0	66.0	23.0	28.1
4	0	2.288	33	1	137.0	40.0	35.0	43.1
6	3	0.248	26	1	78.0	50.0	32.0	31.0

```
[85]: features=data.drop(["type"],axis=1)
```

```
[86]: scale=StandardScaler().fit(features)
features=scale.transform(features)
```

```
[87]: features_scaled=pd.DataFrame(features,columns=data.columns[[0,1,2,4,5,6,7]])
features_scaled.head()
```

```
[87]:
```

	npreg	ped	age	glu	bp	skin	bmi
0	0.749013	0.352710	1.718098	0.879728	0.043756	0.558557	0.102823
1	-0.756247	-0.444246	-0.052006	-1.168442	-0.446013	-0.014657	-0.918661
2	-0.756247	-0.975549	-0.983640	-1.038400	-0.446013	-0.587871	-0.699772
3	-1.057299	5.148879	0.134321	0.522111	-2.568348	0.558557	1.489123
4	-0.154143	-0.741660	-0.517823	-1.396017	-1.752065	0.271950	-0.276586

## 31 4.2 Build a Model with Multiple K

We constructed our models using the silhouette score method. Silhouette is a technique for interpreting and validating the consistency within clusters of data. We do not know the optimal number of clusters that would yield the most useful results. Therefore, we create clusters by varying K from 2 to 8 and subsequently determine the optimum number of clusters (K) with the assistance of the silhouette score.

```
[88]: import warnings
warnings.filterwarnings("ignore")
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
n_clusters=[2,3,4,5,6,7,8]
for K in n_clusters:
```

```

cluster=KMeans(n_clusters=K,random_state=10)
predict=cluster.fit_predict(features_scaled)
score=silhouette_score(features_scaled,predict,random_state=10)
print("For n_clusters={}, silhouette score is {}".format(K,score))

```

```

For n_clusters=2, silhouette score is 0.23295225438615771
For n_clusters=3, silhouette score is 0.21621845377939813
For n_clusters=4, silhouette score is 0.17184136416592355
For n_clusters=5, silhouette score is 0.17201048775549
For n_clusters=6, silhouette score is 0.16968199563891143
For n_clusters=7, silhouette score is 0.17378025698538754
For n_clusters=8, silhouette score is 0.17441664982213687


```

```

[89]: #Importing KMeans from sklearn

from sklearn.cluster import KMeans
#Now we calculate the Within Cluster Sum of Squared Errors (WSS) for different
    ↪ values of k. Next, we
#choose the k for which WSS first starts to diminish. This value of K gives us
    ↪ the best number of
#clusters to make from the raw data.
wcss=[]
for i in range(1,11):
    km=KMeans(n_clusters=i)
#n_clusters - The number of clusters to form as well as the number of centroids
    ↪ to generate
    km.fit(features_scaled)
    wcss.append(km.inertia_)
#inertia_ -Sum of squared distances of samples to their closest cluster center
    ↪
#The elbow curve
plt.figure(figsize=(12,6))
plt.plot(range(1,11),wcss)
plt.plot(range(1,11),wcss, linewidth=2, color="red", marker ="8")
plt.xlabel("K Value")
plt.xticks(np.arange(1,11,1))
plt.ylabel("WCSS")
plt.show()

```



The optimal value for K is identified by the highest silhouette score. From the above output, it is evident that, for K = 2, the silhouette score is the highest. Consequently, we construct the clusters with K = 2.”

```
[90]: # building a K-Means model for K = 2
model = KMeans(n_clusters= 2, random_state= 10)

# fit the model
model.fit(features_scaled)
```

```
[90]: KMeans(n_clusters=2, random_state=10)
```

Now, let’s explore these two clusters to gain insights about them.

## 32 5. Retrieve the Clusters

```
[92]: data_output =features.copy()
# add a column 'Cluster' in the data giving cluster number corresponding to
↳each observation
data_output['Cluster'] = model.labels_
# Reset the index, starting from 1
data_output.index = range(1, len(data_output) + 1)

# head() to display top five rows
data_output.head()
```

**IndexError**

Traceback (most recent call last)

```

Cell In[92], line 3
      1 data_output =features.copy()
      2 # add a column 'Cluster' in the data giving cluster number corresponding
      ↳to each observation
----> 3 data_output['Cluster'] = model.labels_
      4 # Reset the index, starting from 1
      5 data_output.index = range(1, len(data_output) + 1)

IndexError: only integers, slices (':'), ellipsis ('...'), numpy.newaxis ('None'),
↳and integer or boolean arrays are valid indices

```

We have added a column named 'cluster' to the dataframe, indicating the cluster number for each observation.

```

[395]: # 'return_counts = True' gives the number observation in each cluster
np.unique(model.labels_, return_counts=True)

```

```

[395]: (array([0, 1]), array([236, 305], dtype=int64))

```

### 33 Plot a barplot to visualize the cluster sizes

```


[396]: # use 'seaborn' library to plot a barplot for cluster size
sns.countplot(data= data_output, x = 'Cluster')

# set the axes and plot labels
# set the font size using 'fontsize'
plt.title('Cluster Sizes', fontsize = 15)
plt.xlabel('Clusters', fontsize = 15)
plt.ylabel('No. of Customers', fontsize = 15)

# add values in the graph
# 'x' and 'y' assigns the position to the text
# 's' represents the text on the plot
plt.text(x = -0.18, y =236, s = np.unique(model.labels_,
↳return_counts=True)[1][0])
plt.text(x = 0.9, y =305, s = np.unique(model.labels_,
↳return_counts=True)[1][1])

plt.show()

```



The second cluster is slightly larger, containing 305 customers.

## 34 Cluster Centers

The cluster centers can give information about the variables belonging to the clusters

```
[401]: # form a dataframe containing cluster centers
# 'cluster_centers_' returns the co-ordinates of a cluster center
centers = pd.DataFrame(model.cluster_centers_, columns= data_output.columns[:
↪7])
# head() to display top five rows
centers.head()
```

```
[401]:      npreg      ped      age      glu      bp      skin      bmi
0  0.607417  0.157349  0.745409  0.497712  0.610658  0.526982  0.499544
1 -0.470001 -0.121752 -0.576776 -0.385115 -0.472509 -0.407763 -0.386533
```

Now, extract the variables in each of the clusters and attempt to assign a name to each cluster based on the variables

## 35 6 Clusters Analysis

6.1 Analysis of Cluster\_1 Here, we analyze the first cluster by: Checking the size of the cluster. Sorting the variables belonging to the cluster. Computing the statistical summary for observations in the cluster.

```
[407]: # sort the variables based on cluster centers
cluster_1 = sorted(zip(list(centers.iloc[0,:]), list(centers.columns)), reverse_
↪= True)[:3]
```

```
[408]: # size of a cluster_1
np.unique(model.labels_, return_counts=True)[1][0]
```

```
[408]: 236
```

```
[409]: # retrieve the top 3 variables present in the cluster
cluster1_var = pd.DataFrame(cluster_1)[1]
cluster1_var
```

```
[409]: 0    age
      1    bp
      2  npreg
      Name: 1, dtype: object
```

Here, we conduct an analysis of the first cluster, initially examining its size, followed by sorting the variables that belong to the cluster. Subsequently, we compute a statistical summary for the observations within the cluster.

Upon inspection, the first cluster comprises 236 observations. The top three variables in this cluster, ranked by importance, are age, blood pressure (bp), and number of pregnancies (npreg). This suggests that these factors play a significant role within the cluster and may warrant further investigation or attention in the context of the overall dataset.

```
[415]: # get summary for observations in the cluster
# consider the number of orders and customer gender for cluster analysis
data_output[['age', 'bp', 'bmi', 'glu', 'ped']][data_output.Cluster == 0].
describe()
```

```
[415]:
```

	age	bp	bmi	glu	ped
count	236.000000	236.000000	236.000000	236.000000	236.000000
mean	39.559322	78.944915	36.318644	136.249494	0.559343
std	11.152744	10.126074	6.703330	32.529391	0.381184
min	21.000000	50.000000	19.600000	57.000000	0.085000
25%	30.000000	72.000000	32.400000	111.750000	0.263000
50%	39.000000	78.000000	35.550000	135.500000	0.447500
75%	46.000000	85.000000	39.825000	160.250000	0.728000
max	81.000000	110.000000	67.100000	199.000000	2.420000

## 36 6.2 Analysis of Cluster\_\_2

Here, we analyze the second cluster by: Checking the size of the cluster. Sorting the variables belonging to the cluster. Computing the statistical summary for observations in the cluster.

```
[413]: # sort the variables based on cluster centers
cluster_2 = sorted(zip(list(centers.iloc[1,:]), list(centers.columns)), reverse_
↳ True)[:3]
```

```
# size of a cluster_2
np.unique(model.labels_, return_counts=True)[1][1]

# retrieve the top 10 variables present in the cluster
cluster2_var = pd.DataFrame(cluster_2)[1]
cluster2_var
```

```
[413]: 0    ped
      1    glu
      2    bmi
      Name: 1, dtype: object
```

```
[414]: # get summary for observations in the cluster
      # consider the number of orders and customer gender for cluster analysis
data_output[['age', 'bp', 'bmi', 'glu', "ped"]][data_output.Cluster == 1].
      ↪describe()
```

```
[414]:
```

	age	bp	bmi	glu	ped
count	305.000000	305.000000	305.000000	305.000000	305.000000
mean	25.367213	65.675410	30.246557	109.094495	0.462685
std	4.553365	10.523918	5.718282	23.302642	0.311494
min	21.000000	24.000000	18.200000	56.000000	0.085000
25%	22.000000	60.000000	25.900000	93.000000	0.249000
50%	24.000000	66.000000	29.900000	106.000000	0.400000
75%	27.000000	72.000000	34.200000	122.000000	0.583000
max	44.000000	90.000000	55.000000	193.000000	2.288000

37 It can be observed that, in the second cluster, most women exhibit lower mean values for features such as ped, age, bp, bmi, and glucose compared to the first cluster. Higher values in these features are often associated with diabetes. Therefore, we may categorize the first cluster as the ‘diabetes group’ and the second cluster as the ‘non-diabetes group,’ suggesting potential differences in health characteristics between the two clusters.

```
[438]: from sklearn.metrics import accuracy_score
import pandas as pd

# Assuming df is your DataFrame
data_1['type'] = pd.to_numeric(data_1['type'], errors='coerce')
# Replace 0 with 1 and 1 with 0
data_output.replace({0: 1, 1: 0}, inplace=True)
# Drop rows with NaN values, if any
```

```
#df = df.dropna()

# Now you can calculate accuracy
accuracy = accuracy_score(data_output['Cluster'], data_1['type'])
print("Accuracy:", accuracy)
```

Accuracy: 0.711645101663586

```
[439]: data_output.head()
```

```
[439]:
```

	npreg	ped	age	glu	bp	skin	bmi	Cluster
1	6	0.627	50	148.0	72.0	35.0	33.6	1
2	0	0.351	31	85.0	66.0	29.0	26.6	0
3	0	0.167	21	89.0	66.0	23.0	28.1	0
4	1	2.288	33	137.0	40.0	35.0	43.1	0
5	3	0.248	26	78.0	50.0	32.0	31.0	0

In this data frame, '1' represents non-diabetes, and '0' represents diabetes. These labels were assigned through cluster analysis. However, we have the actual labels available, allowing us to compare them with the cluster-assigned labels and calculate the accuracy score.

```
[440]: from sklearn.metrics import roc_auc_score
accuracy = roc_auc_score(data_output['Cluster'], data_1['type'])
print("roc_auc_score:", accuracy)
```

roc\_auc\_score: 0.6934565156988052

```
[444]: from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(data_output['Cluster'], data_1['type'])
cm = confusion_matrix
conf_matrix = pd.DataFrame(data = cm, columns = ['Predicted:0', 'Predicted:1'],
    ↪ index = ['Actual:0', 'Actual:1'])
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap =
    ↪ ListedColormap(['lightskyblue']), cbar = False, linewidths = 0.1, annot_kws =
    ↪ {'size':25})
plt.xticks(fontsize = 20)
plt.yticks(fontsize = 20)
plt.title("Confusion Matrix Based on Clusters")
plt.show()
```



Confusion Matrix Based on Clusters

	Actual:0	255	50
	Actual:1	106	130
		Predicted:0	Predicted:1

The accuracy score is 71.16%, suggesting that the cluster labeling is correct in approximately 71 out of 100 instances. This indicates a reasonably good performance of the cluster labeling method.

[ ]: